

Reality check on Cache Design


- 1 Alternatives to direct mapping
- 2 Enhancing spatial locality

1. Alternative to Direct Mapping

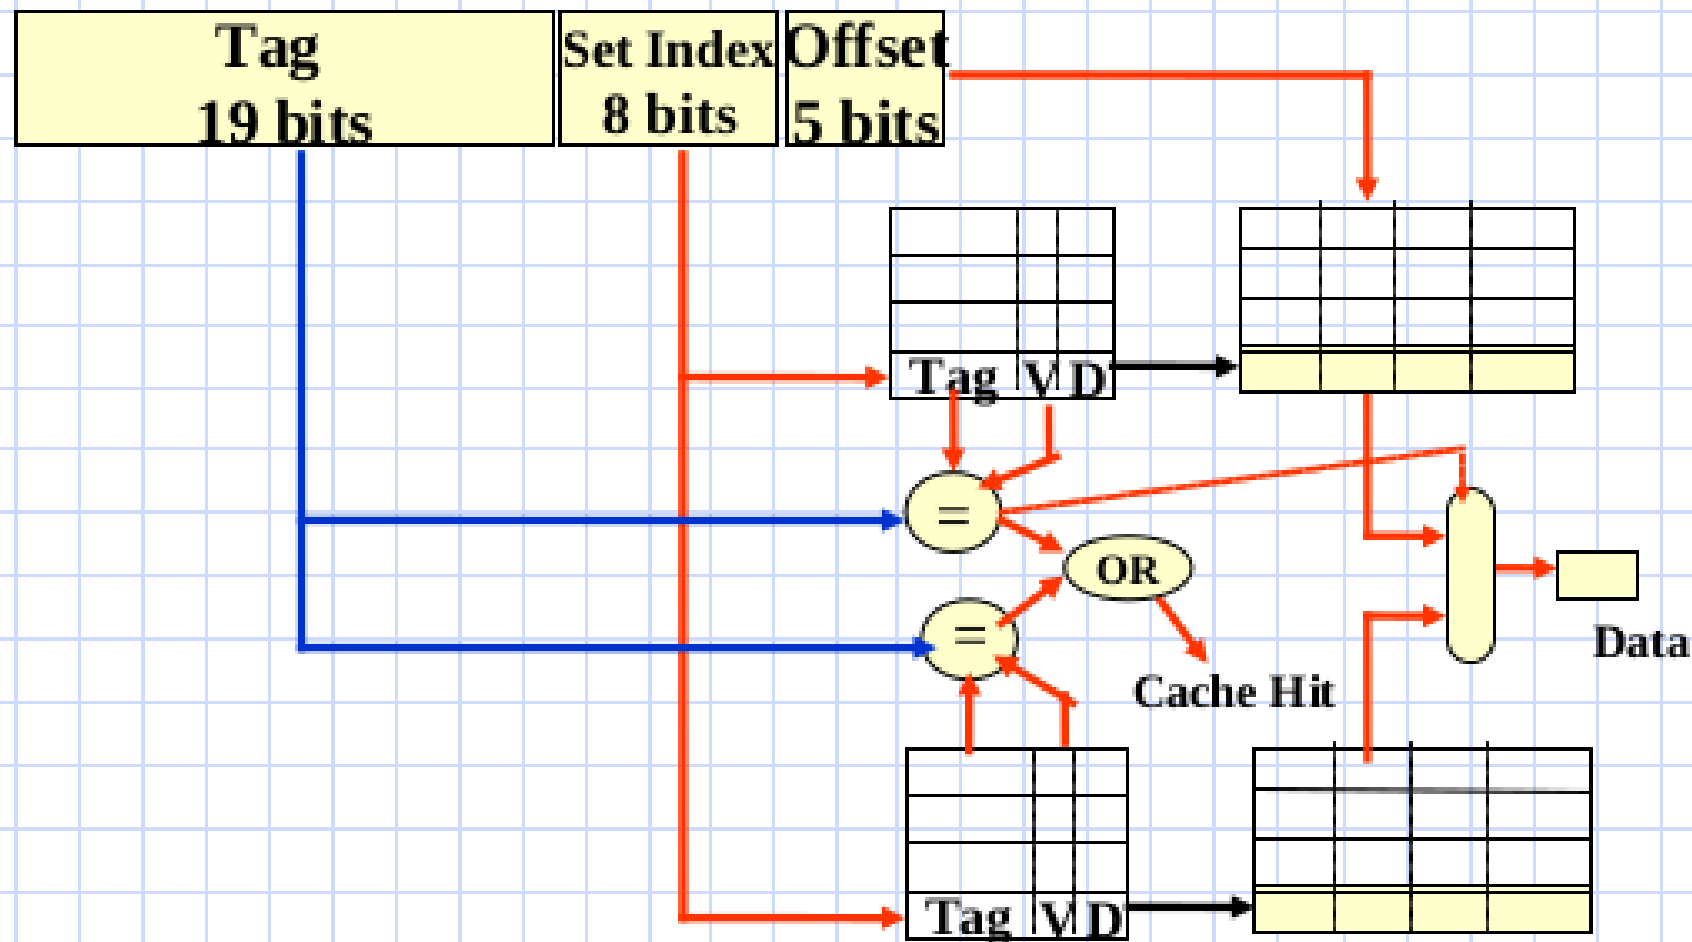
- Why?
 - Conflict misses: Difficult to handle all of them with programming tricks like packing
- Set associative mapping
 - e.g., 2 way set associative
 - Idea: A given memory block can be present in either of 2 blocks of the cache

Set Associative Cache

- A memory block can be loaded into any cache block within a unique set of cache blocks




e.g., 2-way Set Associative Cache



e.g., 4-way Set Associative Cache

- Assume 16KB cache, 32B block size
- $16\text{KB}/32\text{B} = 512$ blocks
- $512/4 = 128$ sets of blocks
- $\log_2 128 = 7$ set index bits
- $\log_2 32 = 5$ offset bits



2. Enhancing Spatial Locality

- Recall: Our prefetching loop for Ex 1
- Hardware prefetcher: Hardware that initiates prefetch into the cache of data that might be required by the CPU soon
 - e.g., Many current Intel cores have 4 hardware prefetchers, 2 for L1 data cache and 2 for L2 cache
- So, your programs may experience more hits due to spatial locality than we have calculated

Caches and Programming

- Look critically at the “important parts of your program”
 - “Hot” functions, where a significant part of program execution time is spent
 - e.g., In the LINPACK benchmark, 70% of the time is spent in Daxpy()
 - Hot loops within hot functions
- Hot functions can be identified using tools called profilers
 - e.g., gprof in Linux/UNIX systems
 - With gcc: use -pg flag

Measuring cache performance

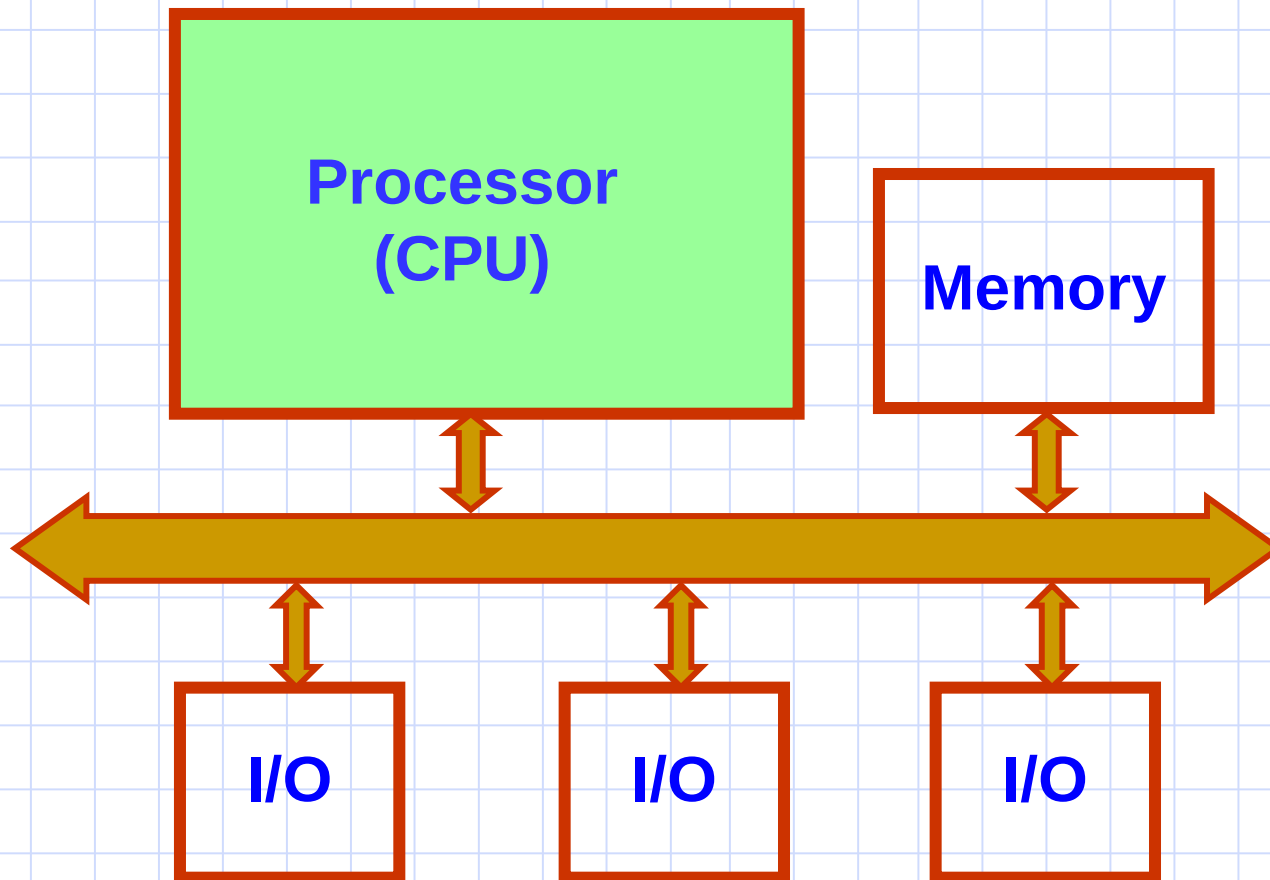
- All modern processors have hardware performance monitoring counters (PMC)
 - They enable counting of numerous “events” during program execution
 - Intel: Separate cycle counter, TSC (time stamp counter), readable with “rdtsc” instruction
 - Other counters require use of privileged instructions

Measuring cache performance

- Related tools
 - VTune amplifier (Intel)
 - PAPI library
 - Perf performance analyzing tool (Linux)
 - You can get privileged access to performance counters on your laptop using “sudo”

```
$ sudo perf stat -e cycles,instructions,cache-references,cache-misses ./a.out
```

Basic Computer Organization



Example: Vector Sum

```
double A[2048], B[2048], C[2048];  
for (i=0; i<2048, i++) C[i] = A[i] + B[i];
```

- What if a CPU has 4 adders?
- It can be designed to support an instruction to do 4 iterations of the Vector Sum loop at a time
 VADD v1_A[0:3], v2_B[0:3], v3_C[0:3]
- Called a vector instruction

Multimedia Extensions

- Hardware support for operations on “short vectors” is provided in existing microprocessors
- Example: 256 bit registers, each split into 4x64b (or 8x32b)
 - Maximum vector length
- Example: Intel “x86” processors
 - SSE (Streaming SIMD Extension)
 - AVX (Advanced Vector Extension)

Vectorization of Loops

We will use a generic notation

Instead of

VADD C[0:3], A[0:3], B[0:3]

$C[0:3] = A[0:3] + B[0:3]$

An example of vectorization

- Given maximum vector length, VL

```
for (i=0; i < N; i++)
```

```
    A[i] = A[i] + B[i];
```

```
for (i=0; i < N; i+=VL)
```

```
    A[i:i+VL-1] = A[i:i+VL-1] + B[i:i+VL-1];
```

What if N is not divisible by VL?

```
for (i=0; i < (N - N%VL); i+=VL)
```

```
    A[i:i+VL-1] = A[i:i+VL-1] + B[i:i+VL-1];
```

```
for (; i<N; i++) A[i] = A[i] + B[i];
```

- This technique is called Stripmining

Auto-vectorization with gcc

- Auto-vectorization compiler feature
 - Compiler will analyze loops of your program and try to use vector instructions
- e.g., gcc command line options
 - -ftree-vectorize for autovectorization
 - -fopt-info-vec to get feedback on autovectorization (which loops were vectorized, etc)
 - Note: These require optimization level of at least -O2

Auto-vectorization with gcc

```
gcc -O2 -ftree-vectorize -fopt-info-vec prog.c
```

```
prog.c:18:1 note: loop vectorized
```


Possible complications

- Dependences between statements within the loop

Example 1

```
for (i=0; i < N; i++) {
```

```
    A[i] = B[i] + C[i];
```

```
    D[i] = (A[i] + A[i+1])/2;
```

```
}
```

```
for (i=0; i < (N - N%VL); i+=VL){
```

```
    A[i:i+VL-1] = B[i:i+VL-1] + C[i:i+VL-1];
```


```
    D[i: ... will get wrong value of A[i+1], etc
```

Example 1

```
for (i=0; i < N; i++) {  
    A[i] = B[i] + C[i];  
    D[i] = (A[i] + A[i+1])/2;  
}
```


```
for (i=0; i < N; i++) {  
    temp[i] = A[i+1];  
    A[i] = B[i] + C[i];  
    D[i] = (A[i] + temp[i])/2;  
}
```

- This loop transformation, through copying of data, is called Node Splitting



Example 2

```
for (i=0; i < N; i++) {  
    X = A[i] + 1;  
    B[i] = X + C[i];  
}
```




```
for (i=0; i < N; i++) {  
    temp[i] = A[i] + 1;  
    B[i] = temp[i] + C[i];  
}
```



- Scalar expansion

Example 3

```
for (i=0; i < N; i++) {  
    A[i] = B[i];  
    C[i] = C[i-1] + 1;  
}
```



```
for (i=0; i < N; i++) A[i] = B[i];
```




```
for (i=0; i < N; i++) C[i] = C[i-1] + 1;
```



- Loop fission

Example 4

```
for (j=1; i < N; j++)  
    for (i=2; i < N; i++)  
        A[i,j] = A[i-1, j] + B[i];
```



```
for (i=2; i < N; i++)  
    for (j=1; j<N; j++)  
        A[i,j] = A[i-1,j] + B[i];
```



- Loop interchange