

<http://cds.iisc.ac.in/courses/ds221>

DS221 Introduction to Scalable Systems [3:1]

Algorithms and Data Structures

Instructor: Chirag Jain
(slides from Prof. Simmhan)

©Department of Computational and Data Science, IISc, 2019

This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/)

Copyright for external content used with attribution is retained by their original authors





About this Course

- Designed as a introductory course on systems
 - Computer Architecture, Operating Systems [MJT]
 - Data Structures, Algorithms [CJ]
 - Big Data Systems, Concepts and Programming [CJ]
 - Parallel Concepts and Programming [VSS]
- Covers **breadth**, not *depth*. Precursor to:
 - Scalable systems for data science, DS256
 - Parallel programming, DS295
 - Algorithmic foundations of big-data biology, DS202



Prerequisites

- Accessible to non-computer science UG Majors
- Requires some prior *programming knowledge*
 - Will NOT teach programming
- We will use C++, Python
 - Makefile, Standard Template Library
 - IDE, GitHub



Sessional

Final

Grading Scheme (CJ)

- Assignments [30 points]
 - 1 Programming Assignments (15 points)
 - 1 Written Assignment (15 points)
- Final Exams [5 points]



Schedule

- Data Structures and Algorithms, Aug 30 – Sept 22
 - Programming Assignment 1
- Big Data Concepts and Spark, Sept 27 – Oct 4
 - Written Assignment (all topics)
- Final Exam, *TBD*
- *Tutorial sessions will be announced*



Ethics

- **IISc POLICY FOR ACADEMIC INTEGRITY**
- Acknowledge and **cite use of others' material**
- **Acknowledges all contributors** to a piece of work
- All work submitted is **his or her own** in a course
- Produce academic work **without the aid** of impermissible materials or collaboration.
- Obtain all results by **ethical means** and report them **accurately**

<https://iisc.ac.in/about/student-corner/academic-integrity/>



Ethics

■ Penalties

- Warning
- Community Service
- Restrictions
- Monetary Penalty
- Withholding Grades
- Suspension
- Expulsion
- Ineligibility



Questions?

<http://cds.iisc.ac.in/courses/ds221>

Data Structure, Algorithms and Data Systems





Class Resources

■ Website

- <http://cds.iisc.ac.in/courses/ds221/>

■ Textbook

- Data Structures, Algorithms, and Applications in C++, Sartaj Sahni*
 - <http://www.cise.ufl.edu/~sahni/dsaac/>

■ Other resources

- The C++ Programming Language, 3rd Edition, Bjarne Stroustrup
- C++ Standard Template Library, <http://www.cplusplus.com/reference/stl/>
- THE ART OF COMPUTER PROGRAMMING (Volume 1 / Fundamental Algorithms), Donald Knuth
- Introduction to Algorithms, Cormen, Leiserson, Rivest and Stein
- www.geeksforgeeks.org/data-structures/

*<http://www.tatabookhouse.com/data-structures-algorithms-and-applications-in-c-plus-plus--9788173715228?ver=1519159641>

**<http://www.flipkart.com/data-structures-algorithms-applications-c-english-2nd/p/itmeyf6jvka3kzdu>



L1: Introduction



Concepts

- **Algorithm:** Outline, the essence of a computational procedure, with step-by-step instructions
- **Program:** An implementation of an algorithm in some programming language
- **Data structure:** Organisation of data need solve the problem (array, list, hashmap)
- **Algorithmic Analysis:** The expected behaviour of the algorithm you have designed, *before you run it*
- **Empirical Analysis:** The behaviour of the program that implements the algorithm, *by running it*

Why not just run it and see how it behaves?



Limitation of Empirical Analysis

- Need to implement the algorithm
 - Time consuming
- Cannot exhaust all possible inputs
 - Experiments can be done only on a limited to set of inputs
 - May not be indicative of running time for other inputs
- Harder to compare two algorithms
 - Same hardware/environments needs to be used



Limitation of Empirical Analysis

- Example?



How do we design an algorithm?

- Intuition
- Mixture of techniques, design patterns
- Experience (body of knowledge)
- Data structures, analysis

How do we implement a program?

- Preferred High Level Language, e.g. C++, Java, Python
- Map algorithm to language, retaining properties
- Use native data structures, libraries

Then why learn about basic data structures?



Basic Data Structures

Lists



Collections of data

- Data Structures to store **collections of data items of same type**
 - Items also called elements, instances, values...depending on context
- **Primitive types** can be boolean, byte, integer, etc.
- **Complex types** can be user or system defined objects, e.g., node, contact, vertex
- **Operations** on the collection
 - Standard operations to create, modify, access elements
- **Properties** of the collection
 - **Invariants that must be maintained**, irrespective of operations performed
- **Challenge**: Understand how to pick the right data structure for your application!



Collections of data

- Can have different **implementations** for same **abstract** data type
 - All offer **same** operations and invariant guarantees
 - Differ in performance (space/time complexity)
- **Challenge**: Understand how to pick the right implementation!
- Also, do we need to have different copies of code for different data types (e.g., integer, floats)?

Try yourself!

- Learn **templates/generics**. In many collections, the item type does not matter for invariants and operations, and can be replaced by a placeholder type “T”.
- Learn C++ **Standard template library (STL)**. Read up examples of abstract collections and their implementations.
 - <http://www.cplusplus.com/reference/stl/>



Linear List (abstract data type)

■ Properties

- ▶ **Ordered** list of items...precedes, succeeds; first, last
- ▶ **Index** for each item...lookup or address item by index value
- ▶ **Well-defined size** for the list at a point in time...can be empty, size may vary with operations performed
- ▶ Items of **same type** present in the list

■ Operations

- ▶ Create, destroy
- ▶ Add, remove item
- ▶ Lookup by index, item value
- ▶ Find size, check if empty
- ▶ *Precise name of operation may vary with language, but semantics remain same/similar....READ THE DOCS!*

Type = int, Size = 7

<i>Index</i>	0	1	2	3	4	5	6
<i>Item</i>	36	5	75	11	7	19	-1



1-D Array (implementation of list)

- List implementation using arrays, in a prog. language
- Arrays are **contiguous** memory locations with **fixed capacity**
 - Contiguous locations mean **locality** matters!
 - **Capacity** is different from **size**. Size is current number of items in list. Capacity denotes max possible size.
- Allow elements of same type to be present at specific **positions** in the array
 - Position is the **offset** from the start of array memory location, while accounting for **data type size**



Mapping Function

- **Index** in a **List** can be mapped to a **Position** in the **Array**
 - **Mapping function** from index to position
- Say **n** is the capacity of the array
- Simple mapping
 - **position(index) = index**

<u>List Index</u>						
0	1	2	3	4		
					<u>Array Position</u>	
0	1	2	3	4	5	6
36	5	75	11	7	--	--
<u>Item values</u>						n=7



List Operations

- item **get**(index)
- void **set**(index, item)
- int **size**()
- int **capacity**()
- boolean **isEmpty**()
- int **indexOf**(item)



List Operations using Arrays

- void **create**(initCapacity)
 - Create array with initial capacity (*optional hint*)
- void **set**(index, item)
 - Use mapping function to set value at position
 - Sanity checks?
- item **get**(index)
 - Use mapping function to set value at position
 - Sanity checks?



```
class List {           // list with index starting at 1
    int arr[]           // backing array for list
    int capacity        // current capacity of array
    int size            // current occupied size of list

    /**
     * Create an empty list with optional
     * initial capacity provided. Default capacity of 15
     * is used otherwise.
     */
    void create(int _capacity){
        capacity = _capacity > 0 ? _capacity : 15
        arr = new int[capacity]    // create backing array
        size = 0                   // initialize size
    }
}
```




```
// assuming pos = index-1 mapping fn.
void set(int index, int item){
    if(index > capacity) { // grow array, double it at least
        arrNew = int[MAX(index, 2*capacity)]
        // copy all items from old array to new
        // source, target, src start, trgt start, length
        copyAll(arr, arrNew, 0, 0, capacity)
        capacity = MAX(index, 2*capacity) // update var.
        delete(arr) // free up memory
        arr = arrNew
    }
    if(index < 1) {
        cout << "Invalid index:" << index << "Expect >=1"
    } else {
        int pos = index - 1
        arr[pos] = item
        size++
    } // end if
} // end set()
} // end List
```

Try yourself!
Implement **get(index)**



List Operations using Arrays

- **int indexOf(item)**
 - Get “first” index of item with given value
 - Sanity checks?
- **void remove(index)**
 - May replace the item with a NULL or shift all items at (index+1) left by 1



List Operations using Arrays

- Increasing capacity
 - Start with initial capacity given by user, or default
 - When capacity is reached
 - Create array with more capacity, e.g. double it
 - Copy values from old to new array
 - Delete old array space
 - Can also be useful to shrink space
 - Why?
- **Pros & Cons of List using Arrays**



Linked List Representation

■ Problems with array

- Pre-defined capacity, under-usage, cost to move items when full. Fixed-size items (primitives)

■ **Solution:** Grow data structure dynamically when we add or remove ➞ Only use as much memory as required

■ *Linked lists* use **pointers** to contiguous chain items

- **Node** structure contains **item** and pointer to **next** node in List
- Add or remove nodes when setting or getting items

Try yourself!



Print the values of **pointer locations** for array and linked list items. Is there any pattern?



Node & Chain

```
class Node {  
    int item  
    Node* next  
}
```

```
class LinkedList {  
    Node* head  
    int size  
    append() {...}  
    get() {...}  
    set() {...}  
    remove {...}  
}
```







Linked List Operations

head=null


Initial empty list




Add item 6




Add item 4



Add items 8, 2



Remove 8



Remove 6




Matrices & n-D Arrays



Matrices & n-D Arrays

- Arrays can have more than 1-dimension
 - 2-D Arrays are called **matrices**, higher dimensions called **tensors**
- Arrays have as many **indexes** for access as dimensions
 - $A[i]$, $B[i][j]$, $C[i][j][k]$
- Dimensions may have **different lengths**

	Column 0	Column 1	Column 2	Column 3
Row 0	$a[0][0]$	$a[0][1]$	$a[0][2]$	$a[0][3]$
Row 1	$a[1][0]$	$a[1][1]$	$a[1][2]$	$a[1][3]$
Row 2	$a[2][0]$	$a[2][1]$	$a[2][2]$	$a[2][3]$



- Mapping from **n-D to 1-D array**
 - Items n dimensions “flattened” into **1** dimension
 - Contiguous memory locations in 1-D
 - Native support in programming languages



n-D Arrays as 1-D Arrays

- Convert $A[i][j]$ to $B[k]$... i =row index, j =column index, C =number of cols, R =number of rows
 - **Row Major Order** of indexing: $k=\text{map}(i,j)=i*C+j$
 - **Column Major Order** of indexing: $k=\text{map}(i,j)=j*R+i$
- *How does this look in memory location layout?*
- *How can you extend this to higher dimensions (tensors)?*

0	1	2	3	4	5
6	7	8	9	10	11
12	13	14	15	16	17

(a) Row-major mapping

0	3	6	9	12	15
1	4	7	10	13	16
2	5	8	11	14	17

(b) Column-major mapping

Figure 7.2 Mapping a two-dimensional array



Matrix-Vector Multiplication

Try yourself!

Write optimised codes for matrix-vector multiplication with: (a) matrix stored as 1D array in row-major order, and (b) with column-major order

$$\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} ax + by + cz \\ dx + ey + fz \\ gx + hy + iz \end{bmatrix}$$

Hint: Reading memory in contiguous locations is significantly faster than jumping around among locations.



n-D Arrays

- **Array of Arrays** representation
- First find pointer for row array
- Then lookup value at column offset in row array
- *How does this look in memory location layout?*
- *Pros & cons relative to using 1-D array representation?*




Figure 7.3 Memory structure for a two-dimensional array



Matrix Multiplication

```
// Given  $a[n][n]$ ,  $b[n][n]$ 
```


```
//  $c[n][n]$  initialized to 0
```

```
for ( $i = 0$ ;  $i < N$ ;  $i++$ )
```

```
    for ( $j = 0$ ;  $j < N$ ;  $j++$ )
```

```
        for ( $k = 0$ ;  $k < N$ ;  $k++$ )
```

```
             $c[i][j] += a[i][k] * b[k][j];$ 
```





Sparse Matrices

- Only a small subset of items are populated in matrix
 - Students and courses taken; faculty and courses taught
 - Adjacency matrix of social network graph
 - vertices are people, edges are “friends”
 - Rows and columns are people, cell has 0/1 value
- Why not use regular 2-D matrix?
 - 1-D representation OR
 - Array of arrays representation



Sparse Matrices as 2-D arrays

- Each non-zero item has one entry in list
 - index: $\langle \text{row}, \text{column}, \text{value} \rangle$
 - index is the $(i-1)^{\text{th}}$ non-zero item in row-major order
 - Space taken in **$3 \times \text{NNZ}$** (number of non zero), compared to **$n \times m$** for non-sparse representation

0	0	0	2	0	0	1	0
0	6	0	0	7	0	0	3
0	0	0	9	0	8	0	0
0	4	5	0	0	0	0	0

(a) A 4×8 matrix

terms	0	1	2	3	4	5	6	7	8
row	1	1	2	2	2	3	3	4	4
col	4	7	2	5	8	4	6	2	3
value	2	1	6	7	3	9	8	4	5

(b) Its linear list representation

Figure 7.14 A sparse matrix and its linear list representation



Sparse Matrix Addition

```
while(p < pMax && q < qMax) {           // C is no. of cols in orig. matrix
    p1 = A[p].r*C + A[p].c              // get index for A in orig. matrix
    q1 = B[q].r*C + B[q].c
    if(p1 < q1)                          // Only A has that index
        C[k] = <A[p].r, A[p].c, A[p].val>           // Copy val
        p++
    else if(p1==q1)                      // Both A & B have that index
        C[k] = <A[p].r, A[p].c, A[p].val+B[q].val>   // Add vals
        p++
        q++
    else                                 // Only B has that index
        C[k] = <B[q].r, B[q].c, B[q].val>           // Copy vals
        q++
    k++
}
```



Compressed Sparse Row (CSR)

- Similar to 2-D array, but more **space efficient**
- 3 arrays (*first 2 same as 2-D array representation*)
 - **A[nnz]** stores non-zero values in row-major order
 - **JA[nnz]** stores column index of nnz in A
 - **IA[m+1]** stores cumulative counts of non-zero values
 - $IA[i] = IA[i-1] + \text{number of NNZ in } (i-1)^{\text{th}} \text{ row}$
 - Always, **IA[0] = 0** and **IA[m+1] = NNZ**
 - i^{th} row elements from **A[IA[i]]**. Existence of j^{th} col elements in JA
- Space taken = $2 * \text{NNZ} + (m + 1)$



Compressed Sparse Row (CSR)

- **A[nnz]** non-zero values in row-major order
- **JA[nnz]** column index of nnz in A ... *Column offset* within the row group for a non zero value
- **IA[m+1]** stores cumulative count of non-zero values till (i-1)th row ... Offset to the start of nnz values for the *i*th row in array A

$$\begin{pmatrix} 0 & 0 & 0 & 0 \\ 5 & 8 & 0 & 0 \\ 0 & 0 & 3 & 0 \\ 0 & 6 & 0 & 0 \end{pmatrix}$$

$$\begin{aligned} \mathbf{A} &= [5 \ 8 \ 3 \ 6]_{\text{nnz}} \\ \mathbf{JA} &= [0 \ 1 \ 2 \ 1]_{\text{nnz}} \\ \mathbf{IA} &= [0 \ 0 \ 2 \ 3 \ 4]_{m+1} \end{aligned}$$

Row groups

Column offset

Offset to row group

To access $X[i,j]$:

- Row i values must be present between $A[s]$ and before $A[e]$, where $s = IA[i]$ and $e = IA[i+1]$
- Check corresponding column offset, $JA[s]$ until $JA[e-1]$. If any of these JA values match j , the value is present in the corresponding index of A .
- If $(e-1 < s)$ OR no JA matches, then the value is 0

Try yourself!
Matrix-matrix
addition using CSR



Tasks

■ Self study (Sahni Textbook)

- Chapters 5 & 6 “Linear Lists—Array & Linked Representations”
- Chapter 7, Arrays and Matrices

■ Programming Self Study

- Try out **list** data structure in **C++ STL**
- Define your own **abstract list interface** using **templates/generics** in C++. Implement create, set, get, front and back using a **1-d array** representation.
- Try out **matrix-matrix multiplication, matrix-vector multiplication** in C++



Upcoming tutorials

- **Date TBD:** Tutorial on Makefile, C++ STL
- **Date TBD:** Tutorial on GitHub
- **Date TBD:** Release programming assignment