



DS221

Data Structures, Algorithms & Data Science Platforms

Instructor: Chirag Jain
(slides from Prof. Simmhan)



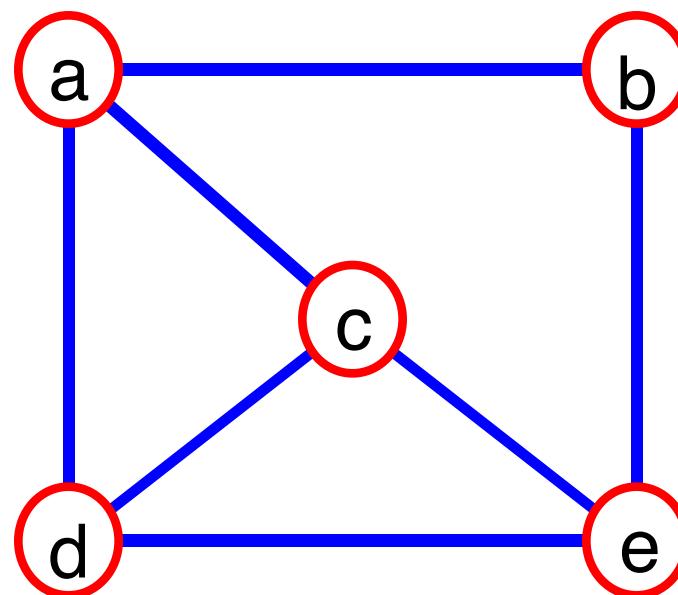
L5: Graphs

Graph ADT, Algorithms



What is a Graph?

- A graph $\mathbf{G} = (\mathbf{V}, \mathbf{E})$ is composed of:
 - \mathbf{V} : set of **vertices**
 - \mathbf{E} : set of **edges** connecting the **vertices** in \mathbf{V}
- An **edge** $e = (u, v)$ is a pair of **vertices**
- Example:



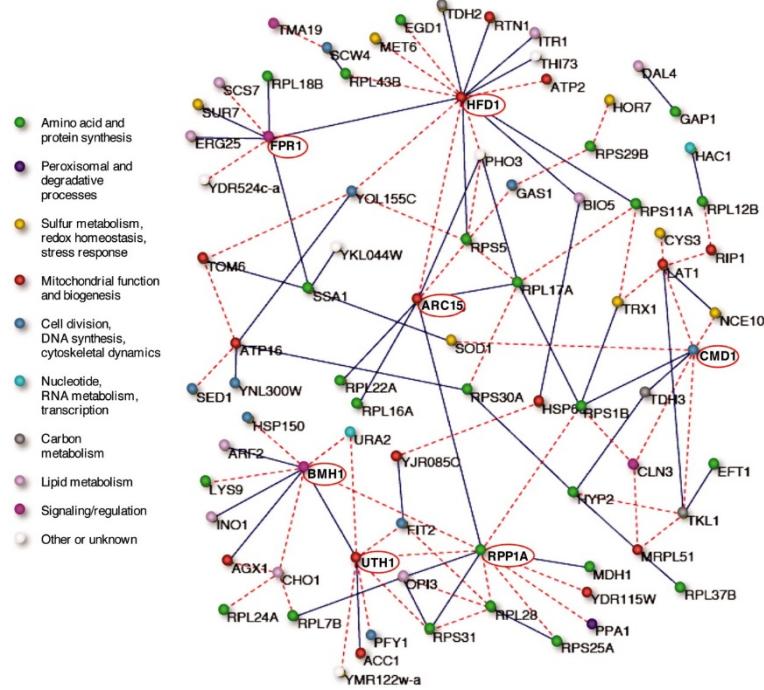
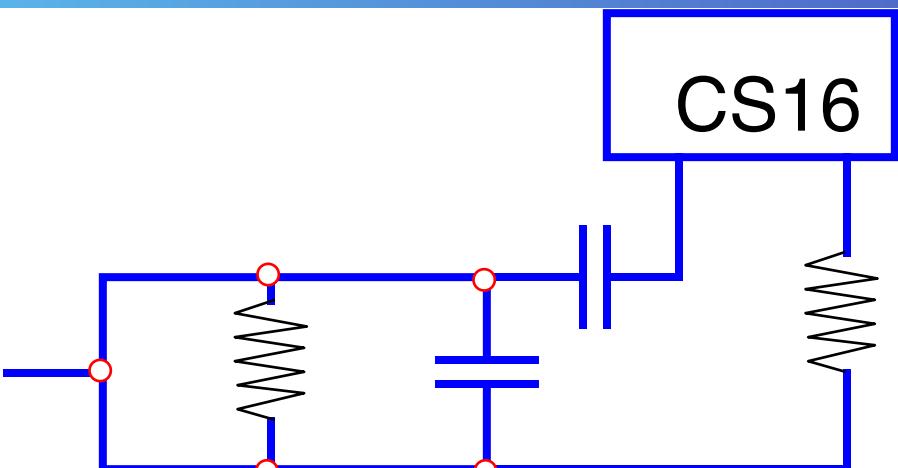
$$\mathbf{V} = \{a, b, c, d, e\}$$

$$\mathbf{E} = \{(a,b), (a,c), (a,d), (b,e), (c,d), (c,e), (d,e)\}$$



Applications

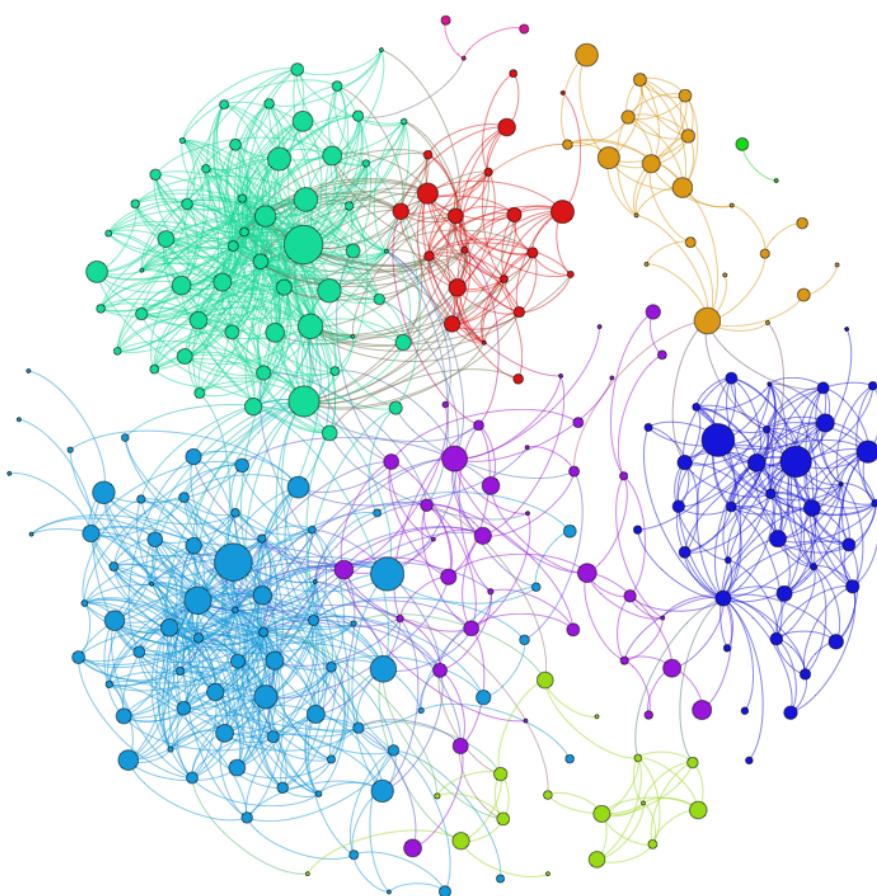
- Electronic circuit design
- Transport networks
- Biological Networks





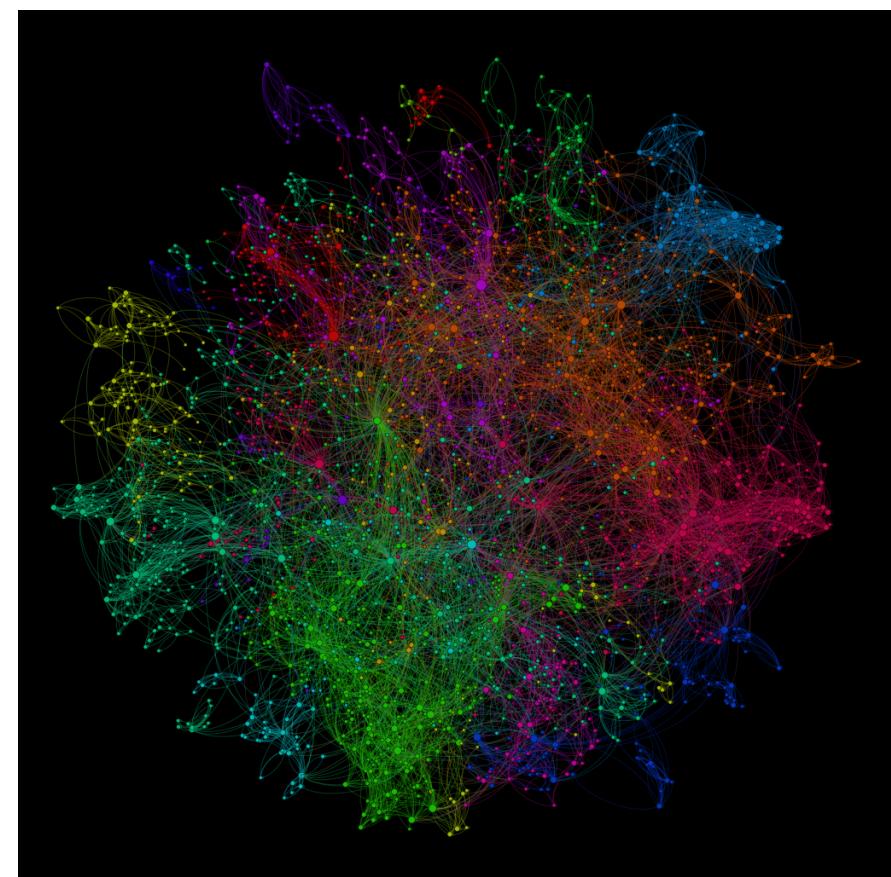
Applications

LinkedIn Social Network Graph



<http://allthingsgraphed.com/2014/10/16/your-linkedin-network/>

Java Call Graph for Neo4J



<http://allthingsgraphed.com/2014/11/12/code-graphs-5-top-open-source-data-projects>



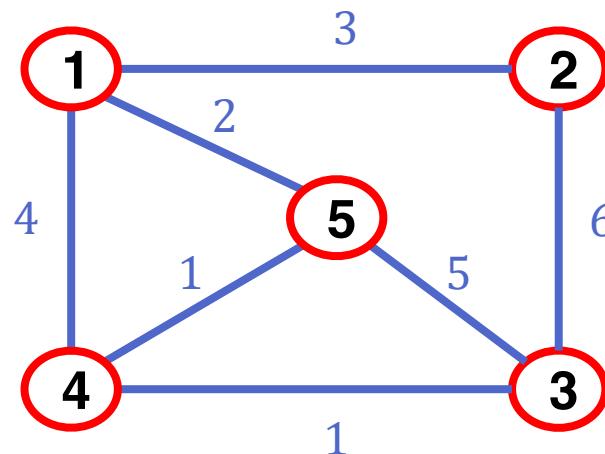
Terminology

- If (v_0, v_1) is an edge in an **undirected** graph,
 - v_0 and v_1 are **adjacent**, or are **neighbours**
 - The edge (v_0, v_1) is **incident** on vertices v_0 and v_1
- If $\langle v_0, v_1 \rangle$ is an edge in a **directed** graph
 - v_0 is the **source vertex** and v_1 is the **target vertex**



Terminology

- Vertices & edges can have **labels** that uniquely identify them
 - Edge label can be formed from the pair of vertex labels it is incident upon...*assuming only one edge can exist between a pair of vertices*
- Edge **weights** indicate some measure of distance or cost to pass through that edge





Terminology

- The **degree** of a vertex is the number of edges incident to that vertex
- For directed graph, also
 - the **in-degree** of a vertex v is the number of edges that have v as the target vertex
 - the **out-degree** of a vertex v is the number of edges that have v as the source vertex
- If d_i is the degree of a vertex i in a graph G with n vertices and e edges, the number of edges is



Terminology

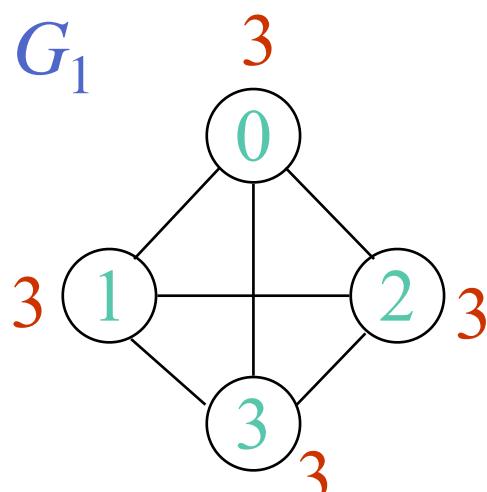
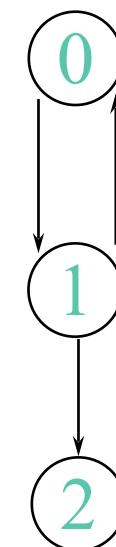
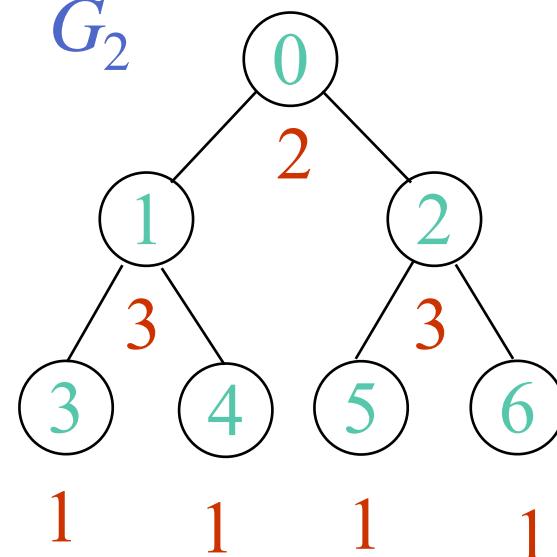
- The **degree** of a vertex is the number of edges incident to that vertex
- For directed graph, also
 - the **in-degree** of a vertex v is the number of edges that have v as the target vertex
 - the **out-degree** of a vertex v is the number of edges that have v as the source vertex
- If d_i is the degree of a vertex i in a graph G with n vertices and e edges, the number of edges is

$$e = \left(\sum_0^{n-1} d_i \right) / 2$$

Why? Since two vertices count an edge, so it will be counted twice



Examples

 G_1  G_2  G_3

in:1, out: 1

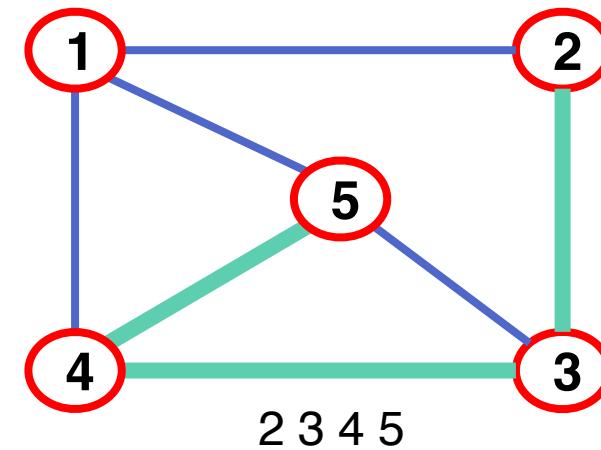
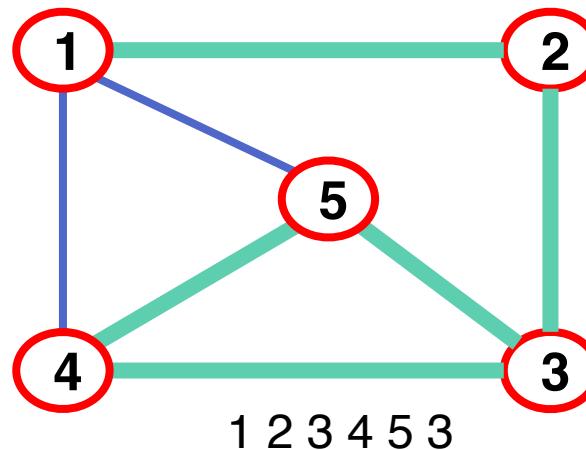
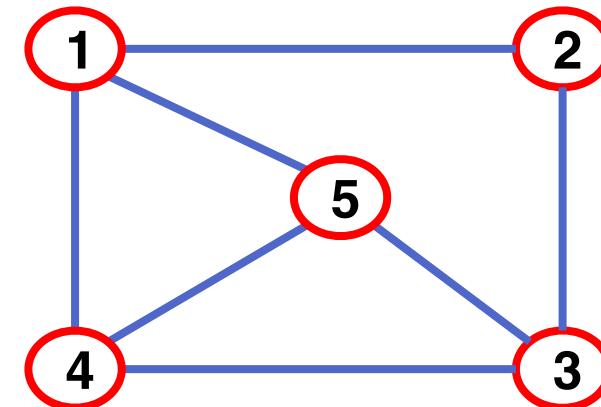
in: 1, out: 2

in: 1, out: 0

*undirected graphs**directed graph*

Terminology

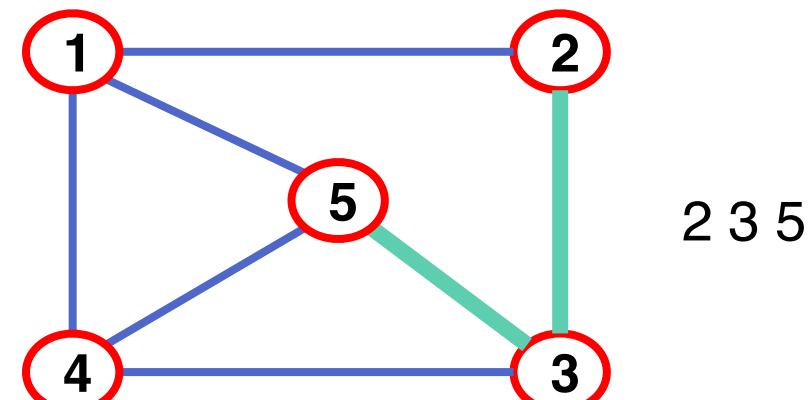
- **path** is a sequence of vertices $\langle v_1, v_2, \dots, v_k \rangle$ such that consecutive vertices v_i and v_{i+1} are adjacent



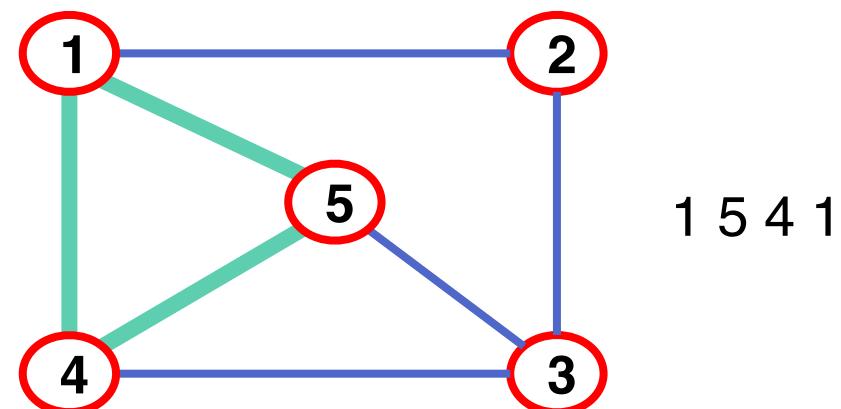


Terminology

- **simple path**: no repeated vertices



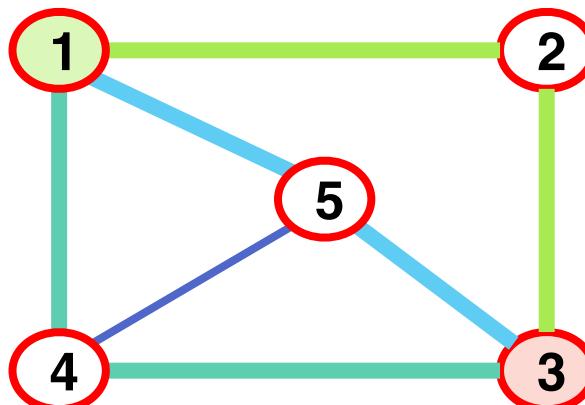
- **cycle**: simple path, except that the last vertex is the same as the first vertex



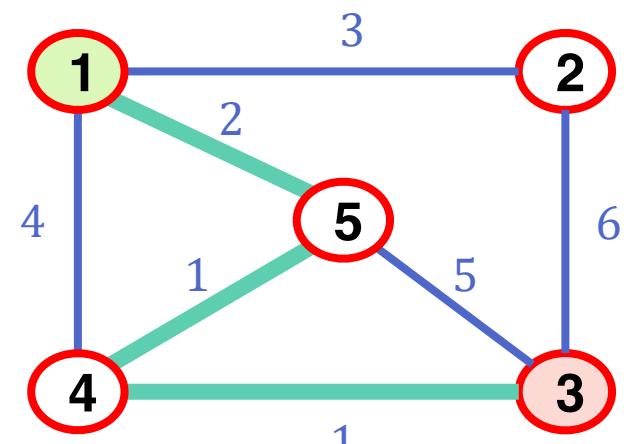


Terminology

- **Shortest Path:** Path between two vertices where sum of the edge weights is the smallest
 - Has to be a simple path (why?)
 - Assume “unit weight” for edges if not specified



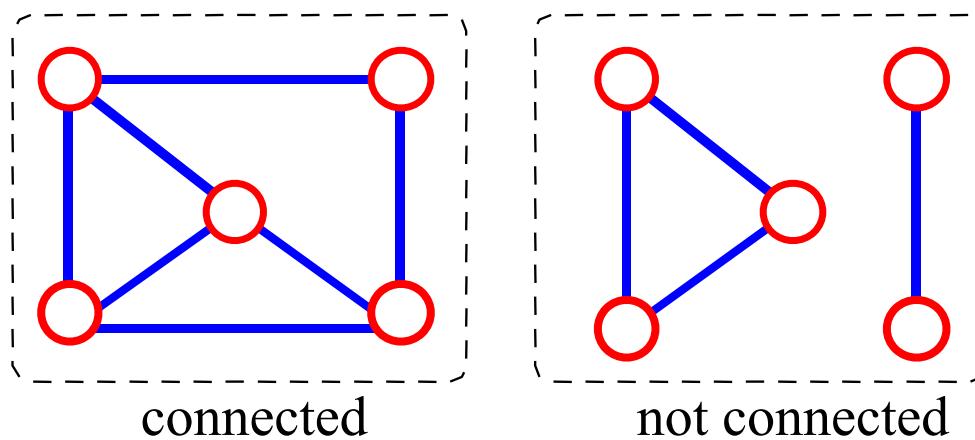
1 2 3
1 5 3
1 4 3



1 5 4 3
13

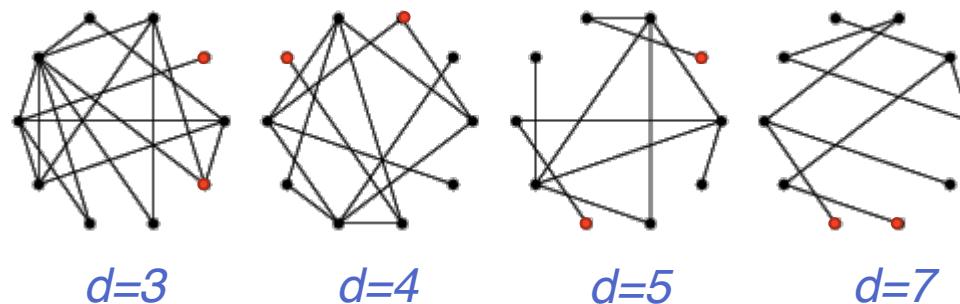
Connected Graph

- **connected graph**: any two vertices are connected by some path



Graph Diameter

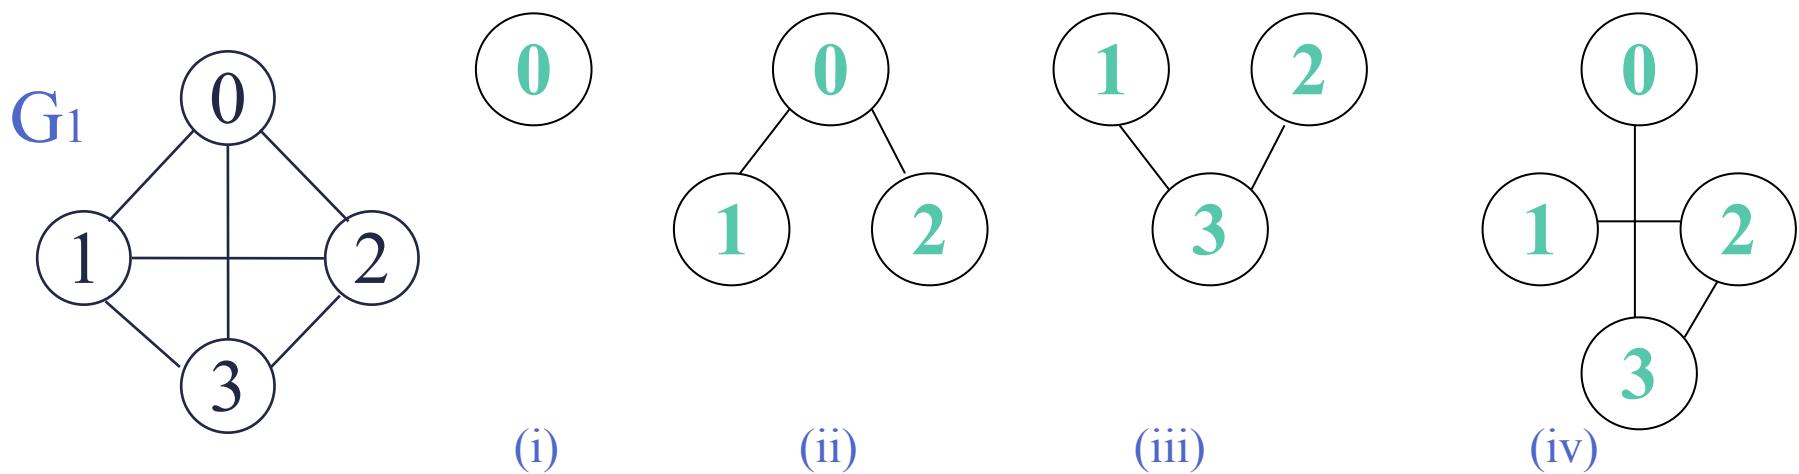
- A **graph's diameter** is the distance of its *longest shortest path*
- if $d(u,v)$ is the distance of the shortest path between vertices u and v , then:
 - $diameter = \text{Max}(d(u,v))$, for all u, v in V
- A disconnected graph has an infinite diameter





Subgraph

- **subgraph**: subset of vertices and edges forming a graph

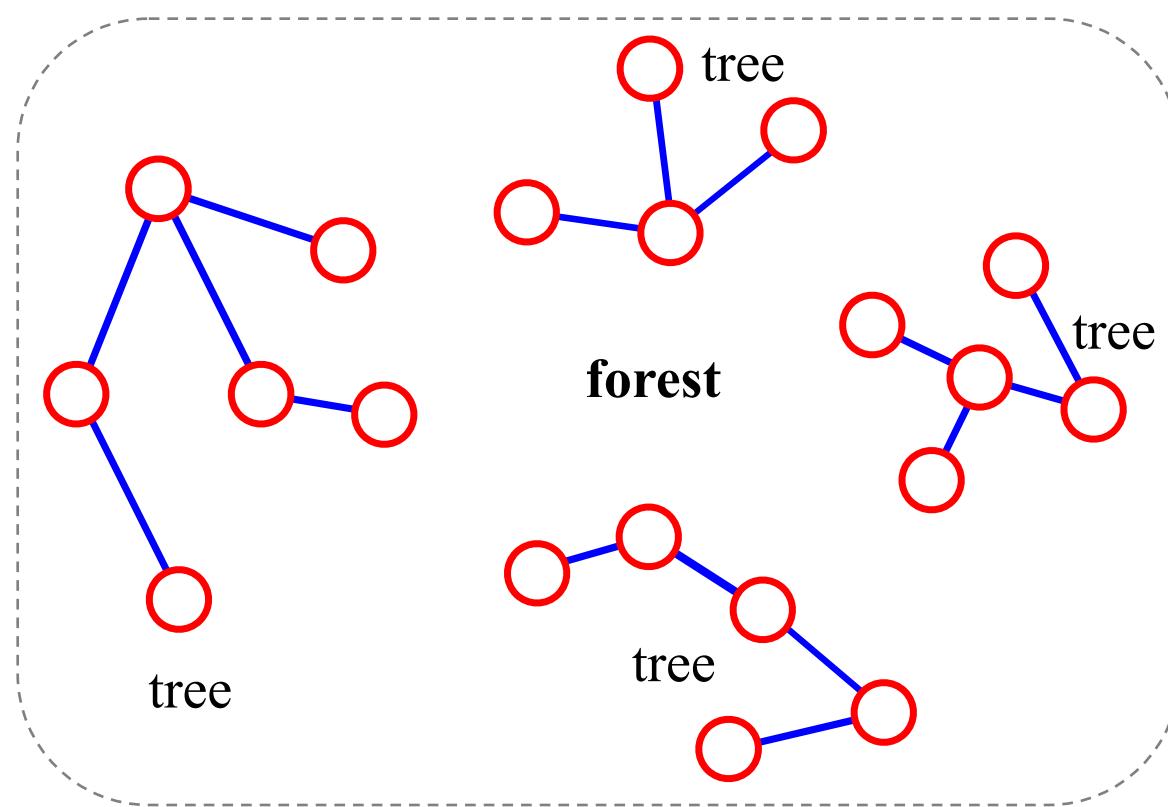


(a) Some of the subgraphs of G_1



Trees & Forests

- **tree** - connected graph without cycles
- **forest** - collection of trees

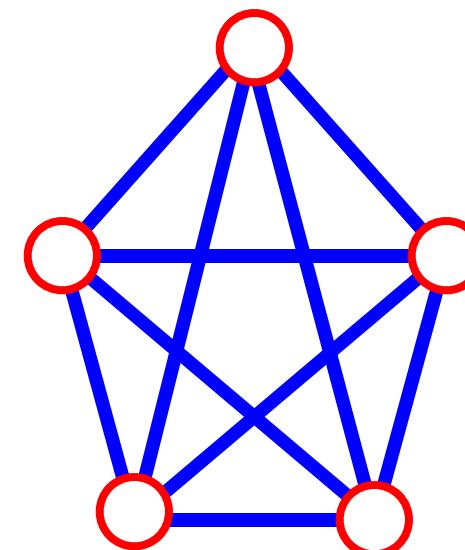




Fully Connected Graph

- Let $n = \# \text{vertices}$, and $m = \# \text{edges}$
- Complete graph (or) Fully connected graph:** One in which all pairs of vertices are adjacent
- Count of edges in a complete graph?*
 - Each of the n vertices is incident to $n-1$ edges, however, we would have counted each edge twice! Therefore, $m = n(n-1)/2$.

If a graph is not complete:
 $m < n(n-1)/2$



$$n = 5$$

$$m = (5*4)/2 = 10$$



More Connectivity

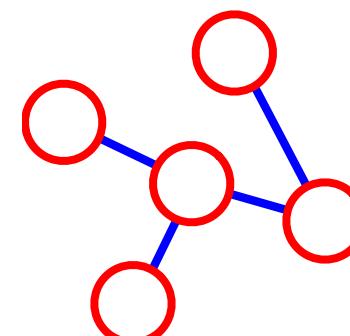
n = #vertices

m = #edges

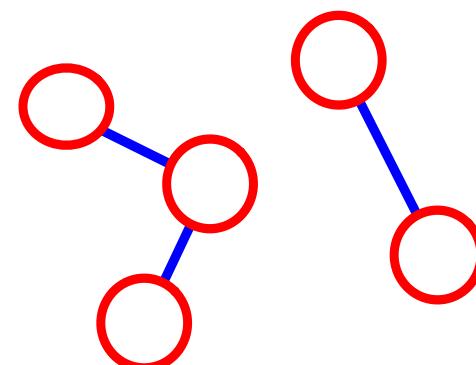
■ For a tree **m** = **n** - 1

(Try to prove using induction)

If **m** < **n** - 1, G is
not connected



n = 5
m = 4

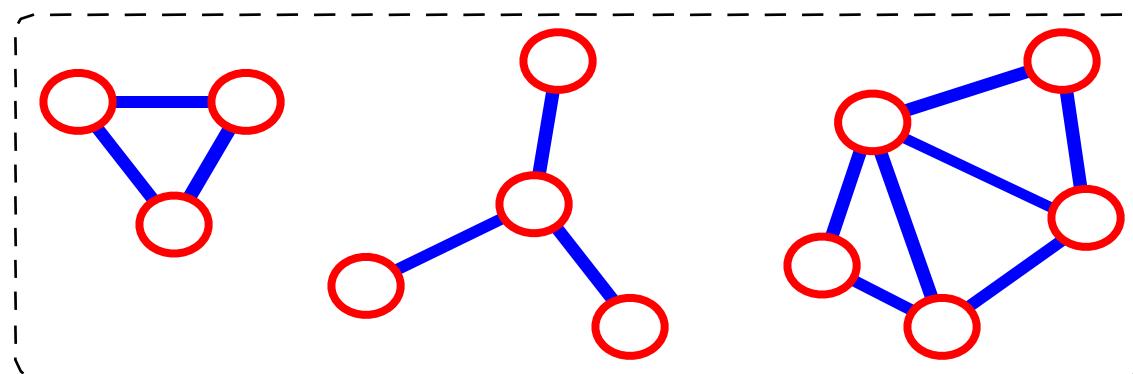


n = 5
m = 3



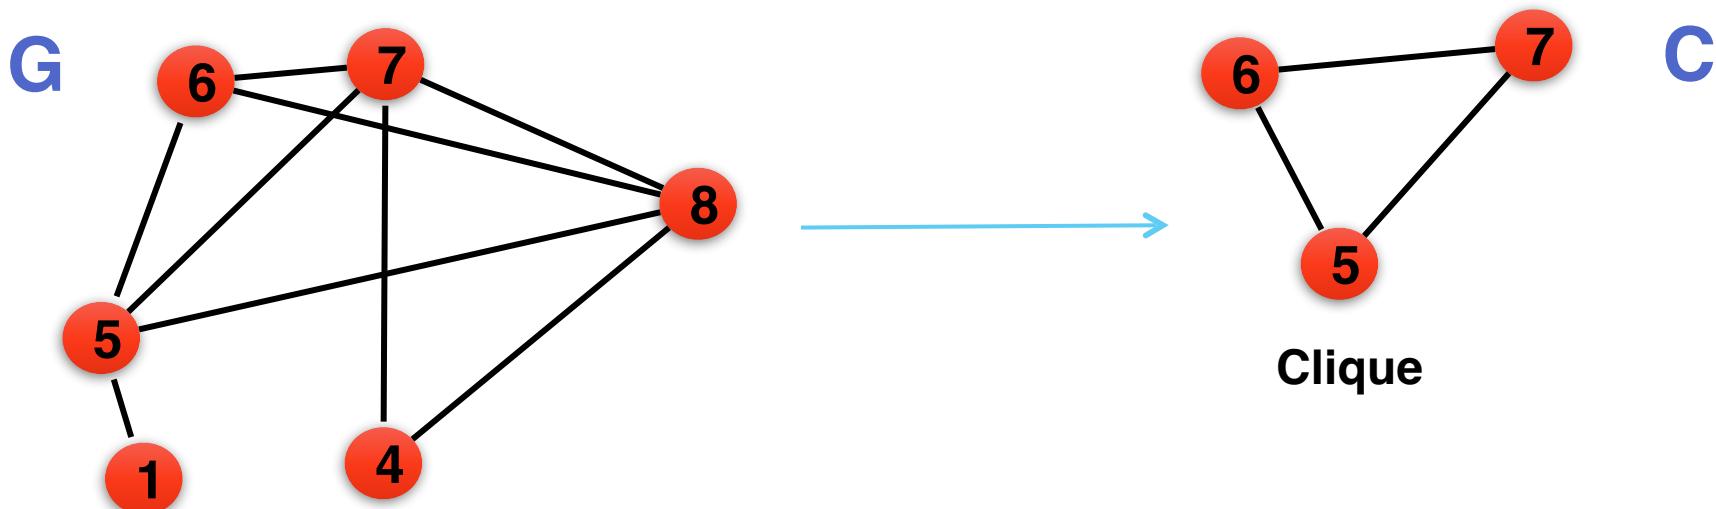
Connected Component

- A **connected component** is a maximal set of nodes such that each pair of nodes is connected by a path
- A connected graph has exactly one connected component.



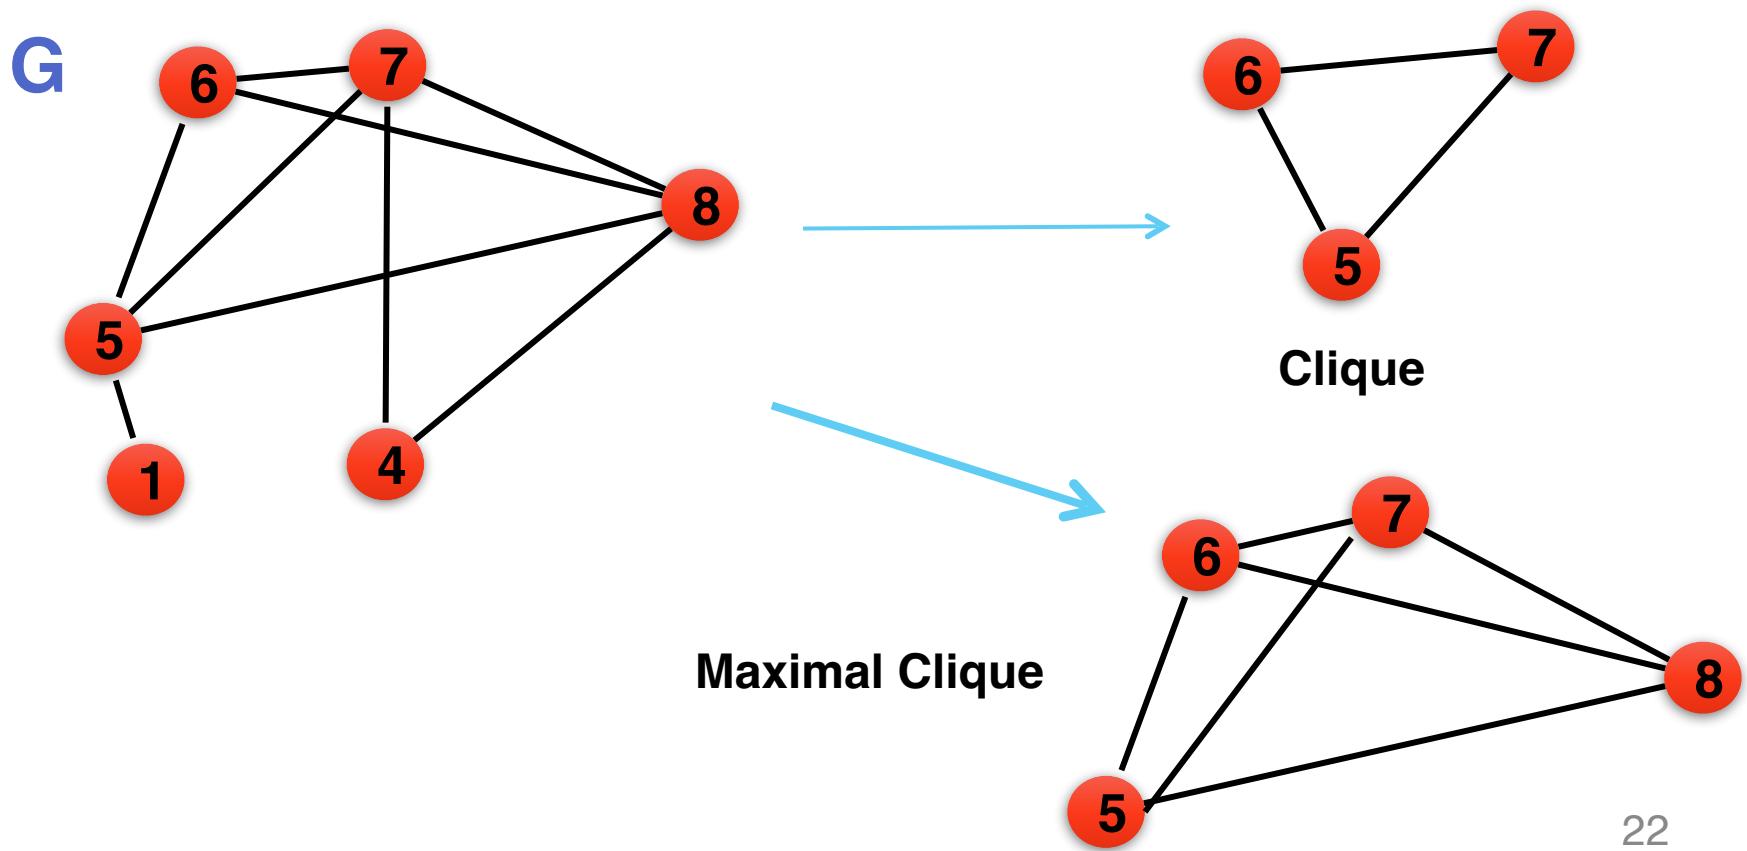
Clique

- A subgraph C of a graph G with *edges between all pairs of vertices*



Maximal Clique

- A maximal clique is a clique that is not part of a larger clique



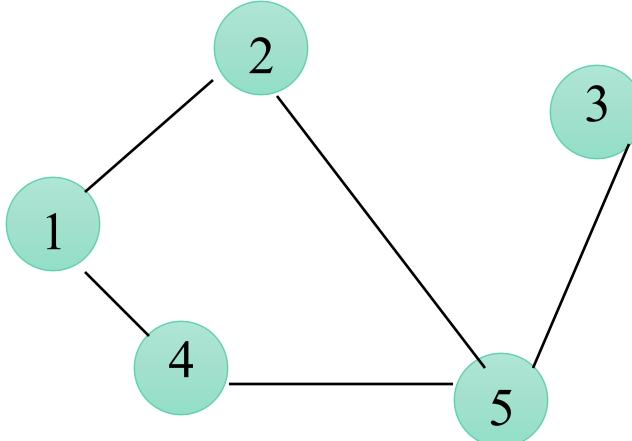


Graph Representation

- Adjacency Matrix
- Adjacency Lists
 - Linked Adjacency Lists
 - Array Adjacency Lists

Adjacency Matrix

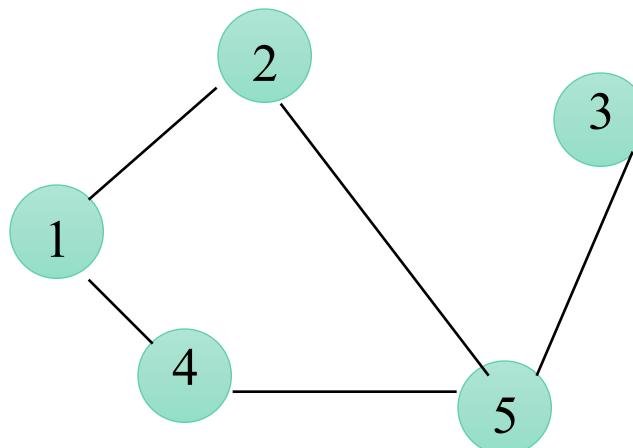
- Boolean $n \times n$ matrix, where $n = \#$ of vertices
- $A(i,j) = 1$ iff (i,j) is an edge



	1	2	3	4	5
1	0	1	0	1	0
2	1	0	0	0	1
3	0	0	0	0	1
4	1	0	0	0	1
5	0	1	1	1	0



Adjacency Matrix Properties

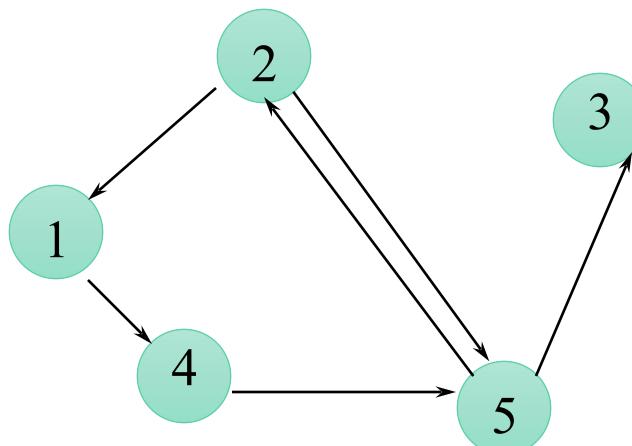


	1	2	3	4	5
1	0	1	0	1	0
2	1	0	0	0	1
3	0	0	0	0	1
4	1	0	0	0	1
5	0	1	1	1	0

- Diagonal entries are zero.
- Adjacency matrix of an *undirected graph* is *symmetric*.
 - $A(i,j) = A(j,i)$ for all i and j .



Adjacency Matrix (Digraph)



	1	2	3	4	5
1	0	0	0	1	0
2	1	0	0	0	1
3	0	0	0	0	0
4	0	0	0	0	1
5	0	1	1	0	0

- Diagonal entries are zero.
- Adjacency matrix of a directed graph need not be symmetric.



Adjacency Matrix

- n^2 bits of space
- For an *undirected graph*, may store only lower or upper triangle (exclude diagonal)
 - $(n^2 - n)/2$ bits
- $O(n)$ time to find vertex degree and/or vertices adjacent to a given vertex.



Adjacency Lists

- Adjacency list for vertex i is a linear list of vertices adjacent from vertex i .
- An array of n adjacency lists.

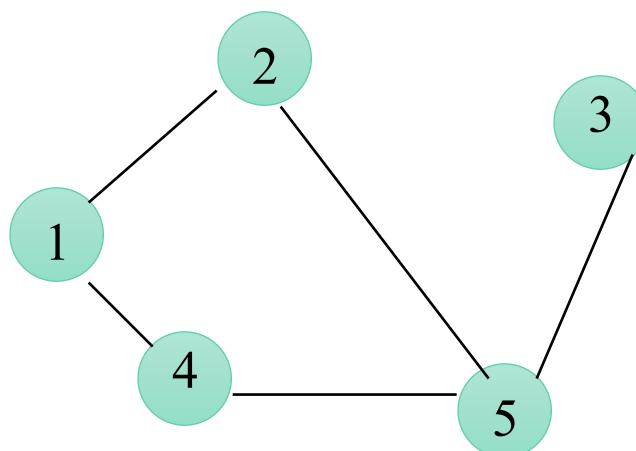
$$\text{aList}[1] = (2,4)$$

$$\text{aList}[2] = (1,5)$$

$$\text{aList}[3] = (5)$$

$$\text{aList}[4] = (5,1)$$

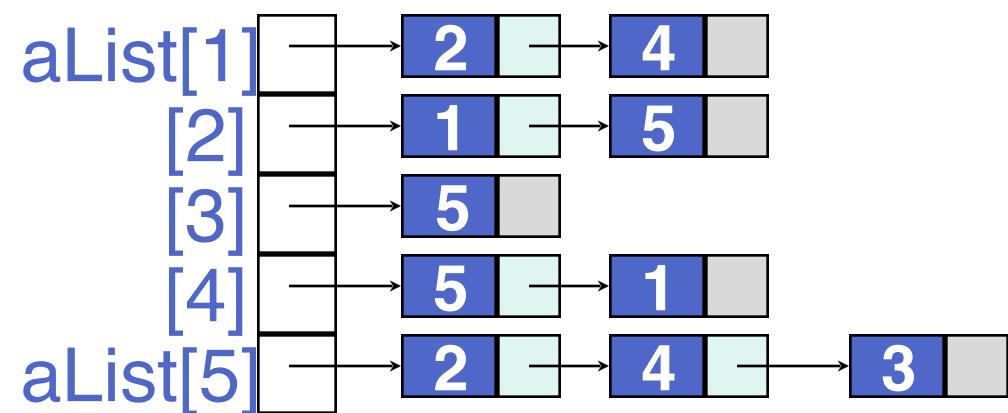
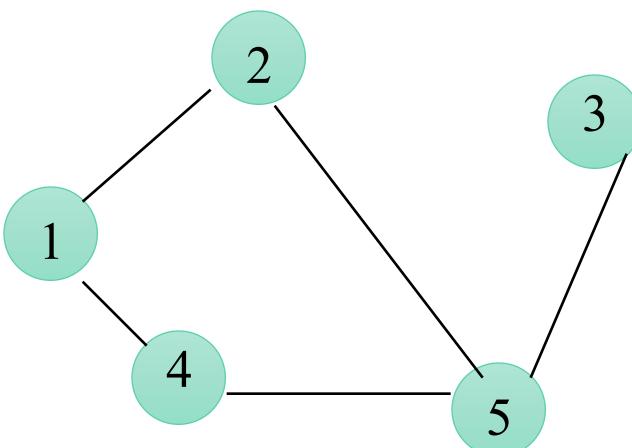
$$\text{aList}[5] = (2,4,3)$$





Linked Adjacency Lists

- Each adjacency list is a chain.



- Array Length = n
- # of chain nodes = $2e$ (undirected graph)
- # of chain nodes = e (digraph)



Storing Weighted Graphs

- Cost adjacency matrix
 - $C(i,j) = \text{cost of edge } (i,j)$ instead of 0/1

- Adjacency lists
 - Each list element is a pair (adjacent vertex, edge weight)



ADT for Graph

```
class Vertex<V,E> {
    int id;
    V value;
    int GetId();
    V GetValue();
    List<Edge<V,E>> Neighbors();
}

class Edge<V,E> {
    int id;
    E value;
    int GetId();
    E GetValue();
    Vertex<V,E> GetSource();
    Vertex<V,E> GetTarget();
}
```



ADT for Graph

```
class Graph<V,E>{
    List<Vertex<V,E>> vertices;
    List<Edge<V,E>> edges;

    void InsertVertex(Vertex<V,E> v);
    void InsertEdge(Edge<V,E> e);

    bool DeleteVertex(int vid);
    bool DeleteEdge(int eid);

    List<Vertex<V,E>> GetVertices();
    List<Edge<V,E>> GetEdges();

    bool IsEmpty(graph);
}
```



Sample Graph Problems

- Graph traversal
 - Searching
 - Shortest Paths
 - Connectedness
- Spanning Tree
- Graph centrality
 - PageRank
 - Betweenness centrality
- Graph clustering
 - K-means clustering

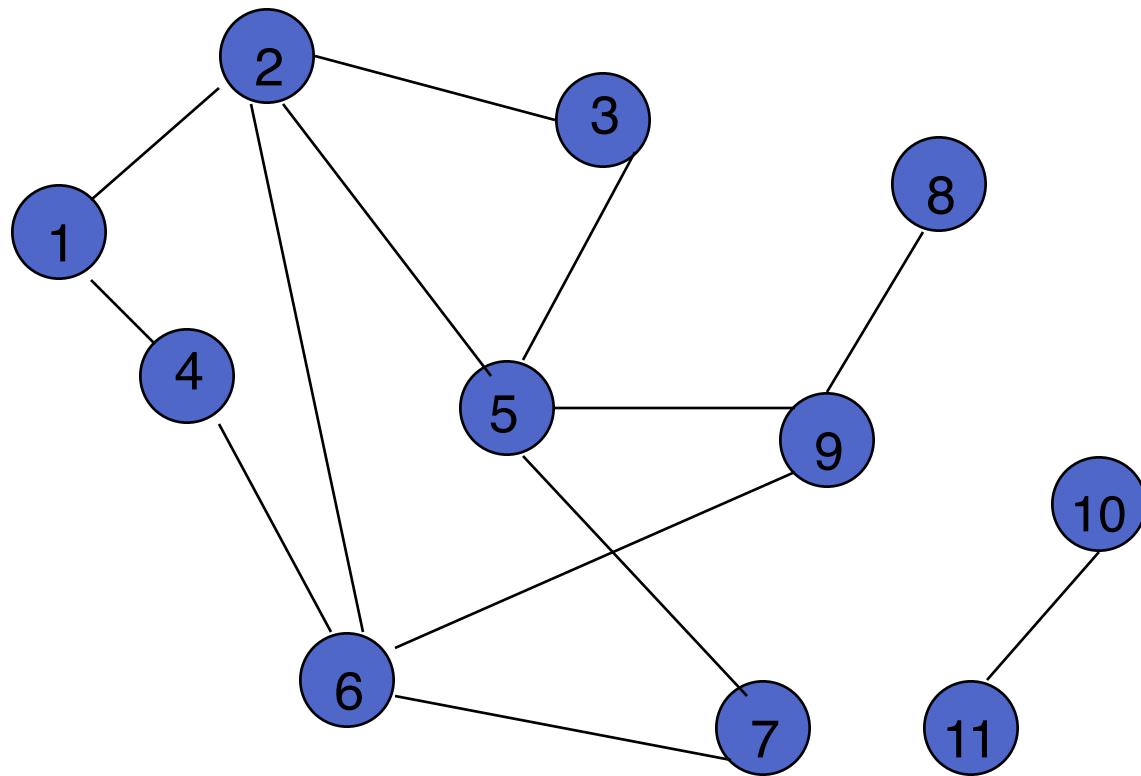


Graph Search & Traversal

- Find a vertex (or edge) with a given ID or value
- Traverse through the graph to list all vertices in a particular order
 - Finding the item can be side-effect of traversal



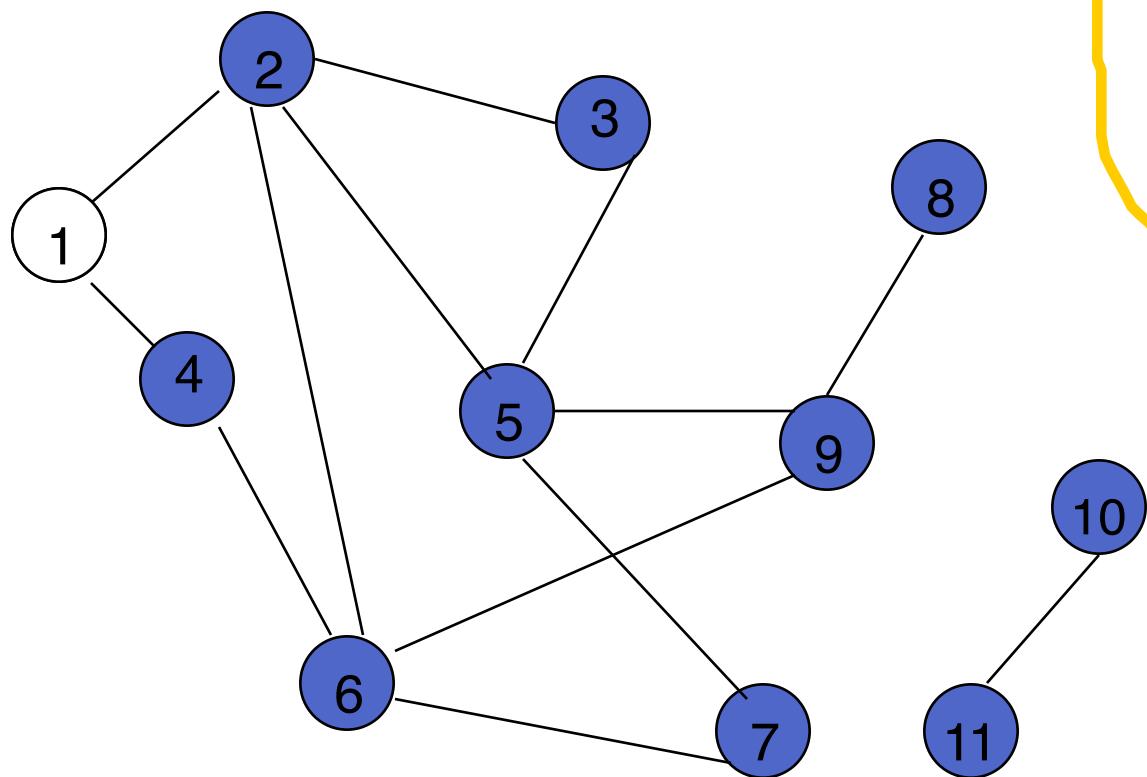
Breadth-First Search Example



Start search at vertex 1.



Breadth-First Search Example



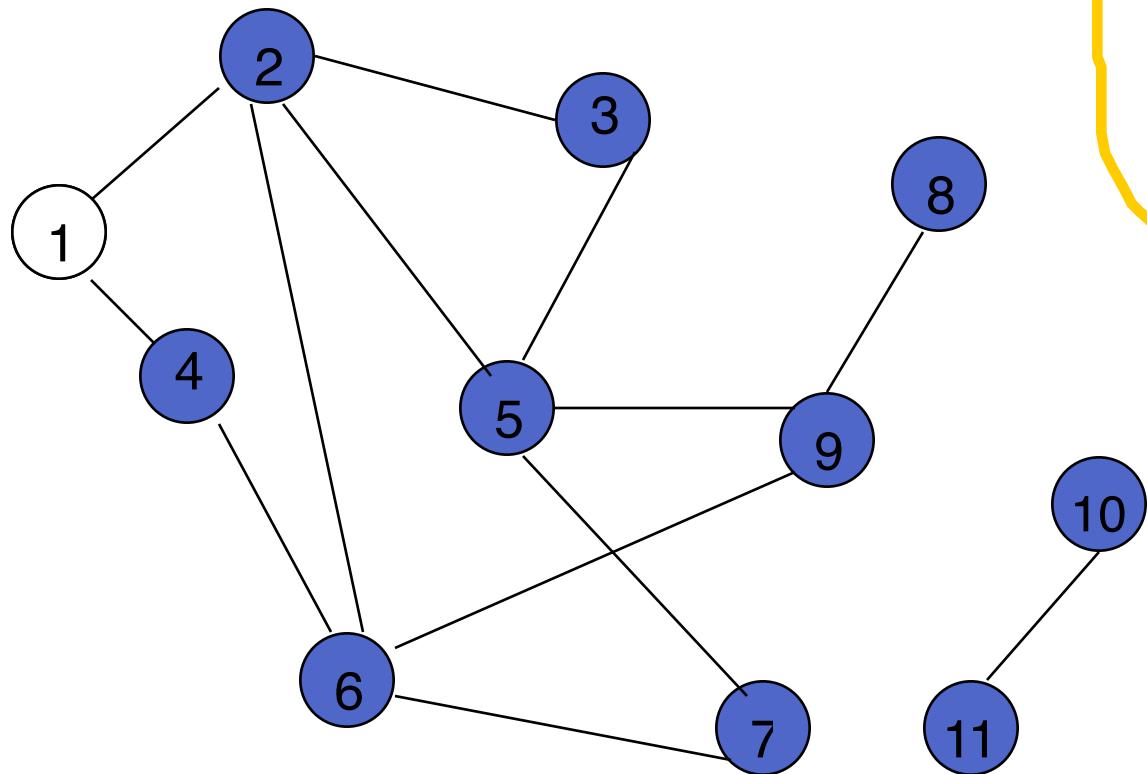
FIFO Queue

1

Visit the start vertex and put in a FIFO queue.



Breadth-First Search Example



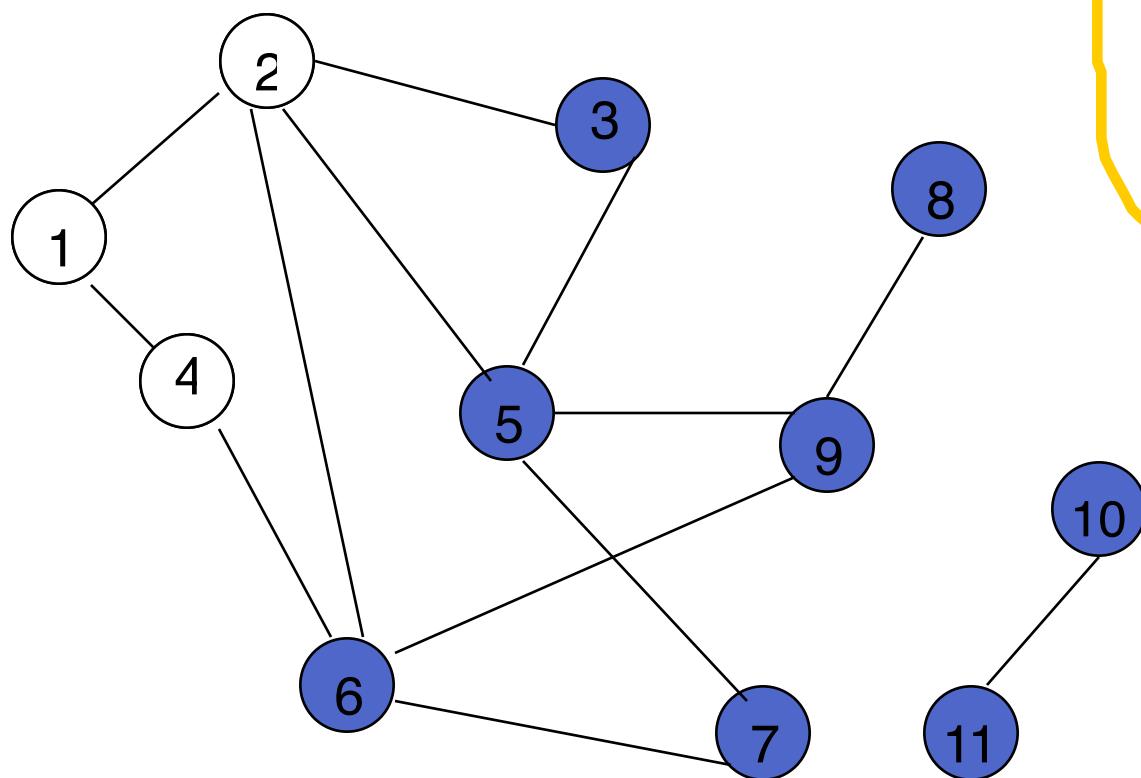
FIFO Queue

1

Remove 1 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



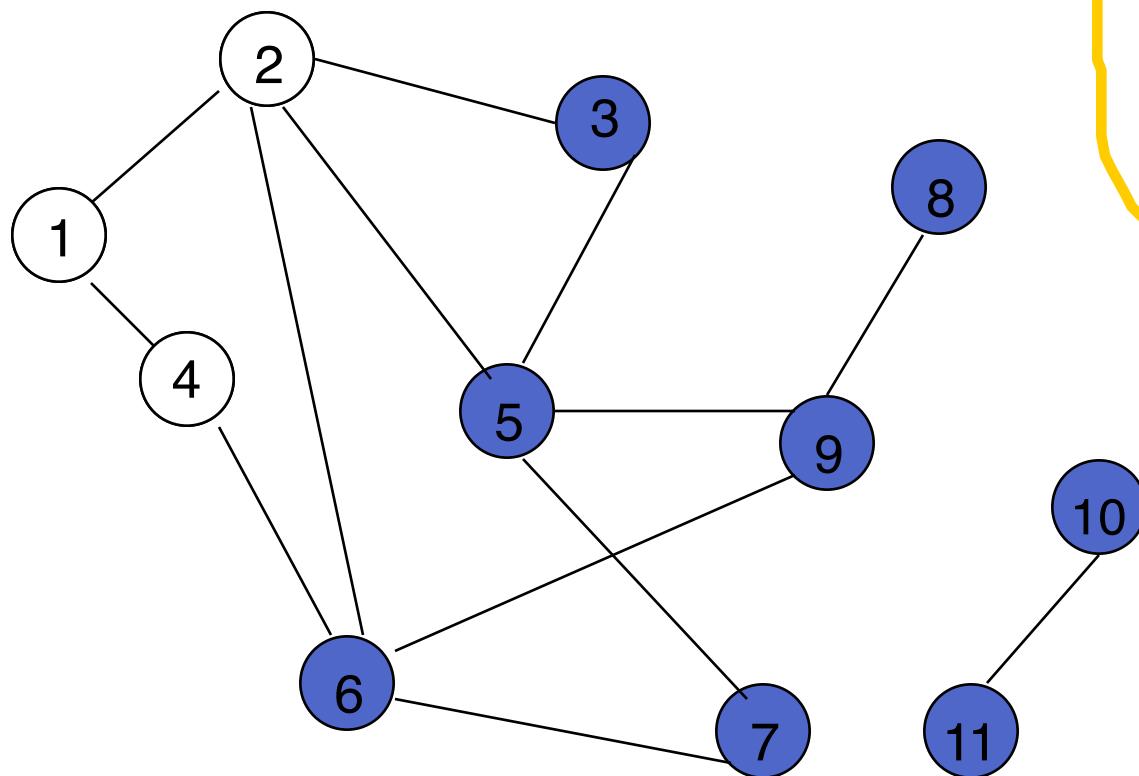
FIFO Queue

2 4

Remove 1 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



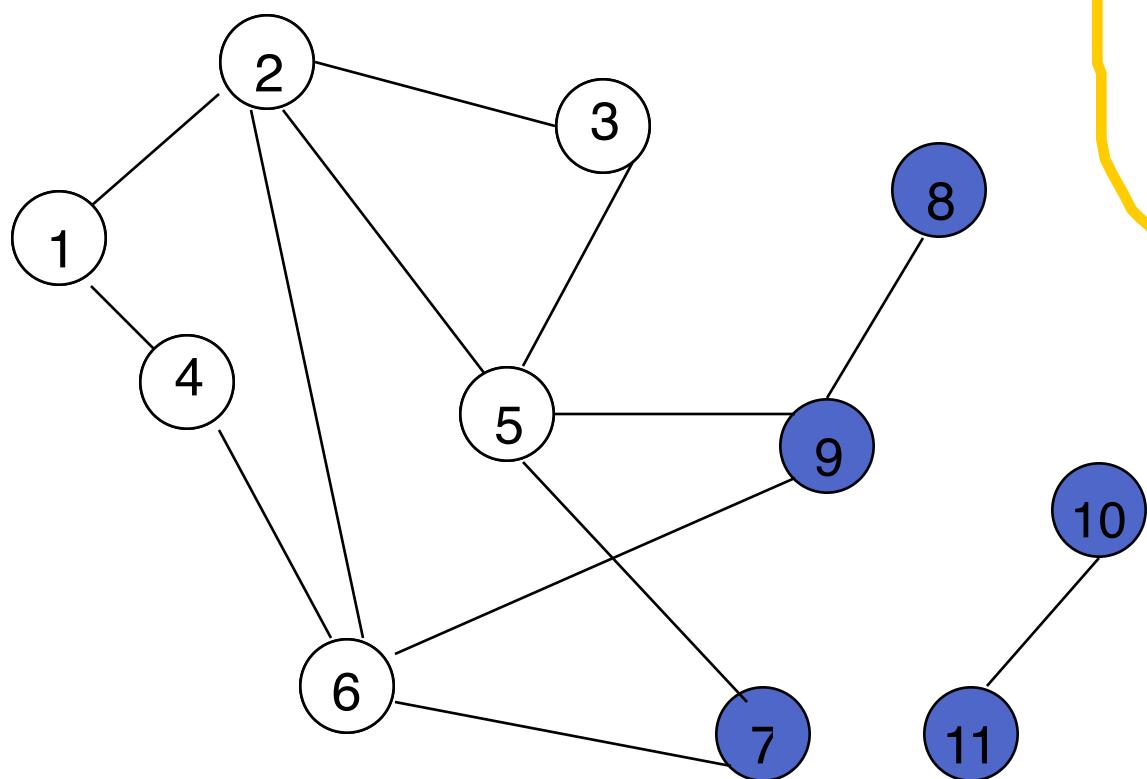
FIFO Queue

2 4

Remove 2 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



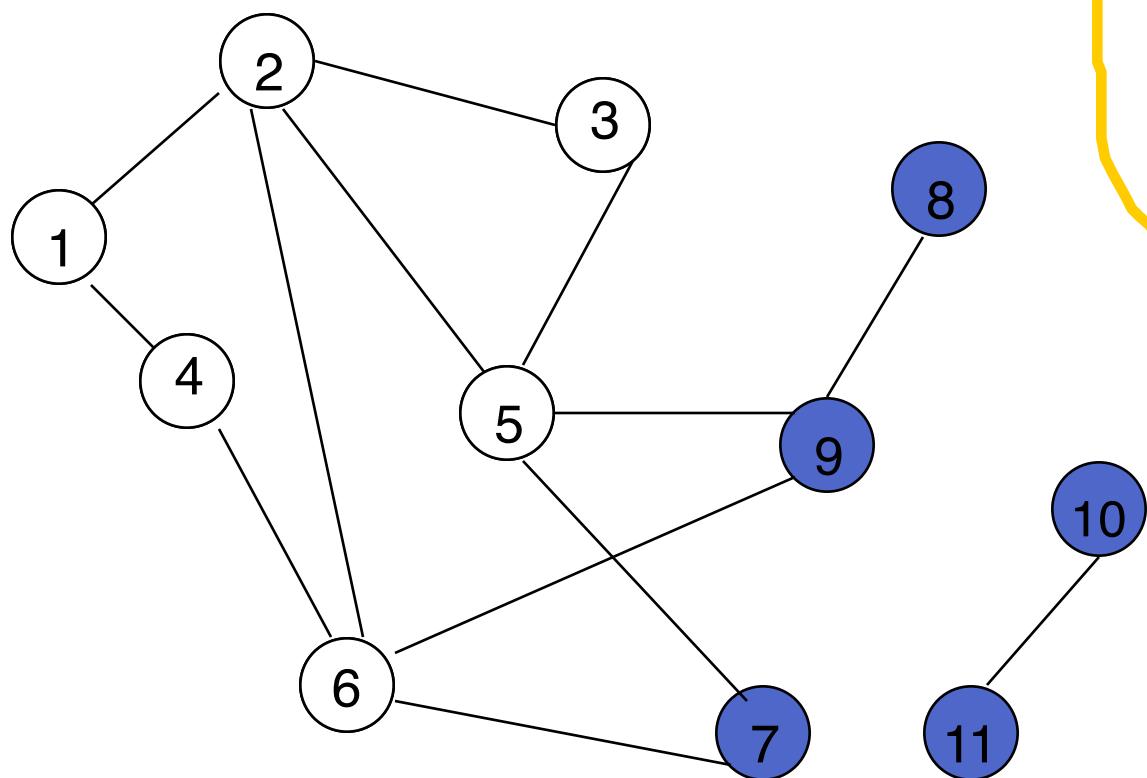
FIFO Queue

4 5 3 6

Remove 2 from Q; visit adjacent unvisited vertices;
put in Q.



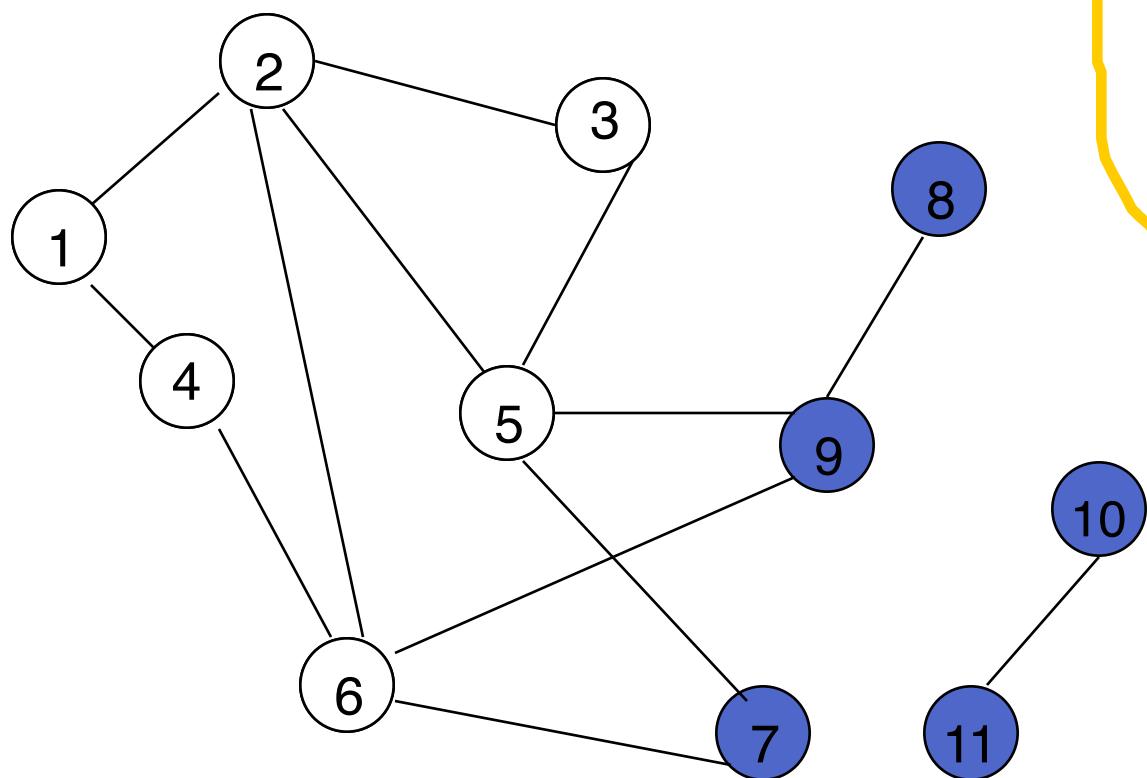
Breadth-First Search Example



Remove 4 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



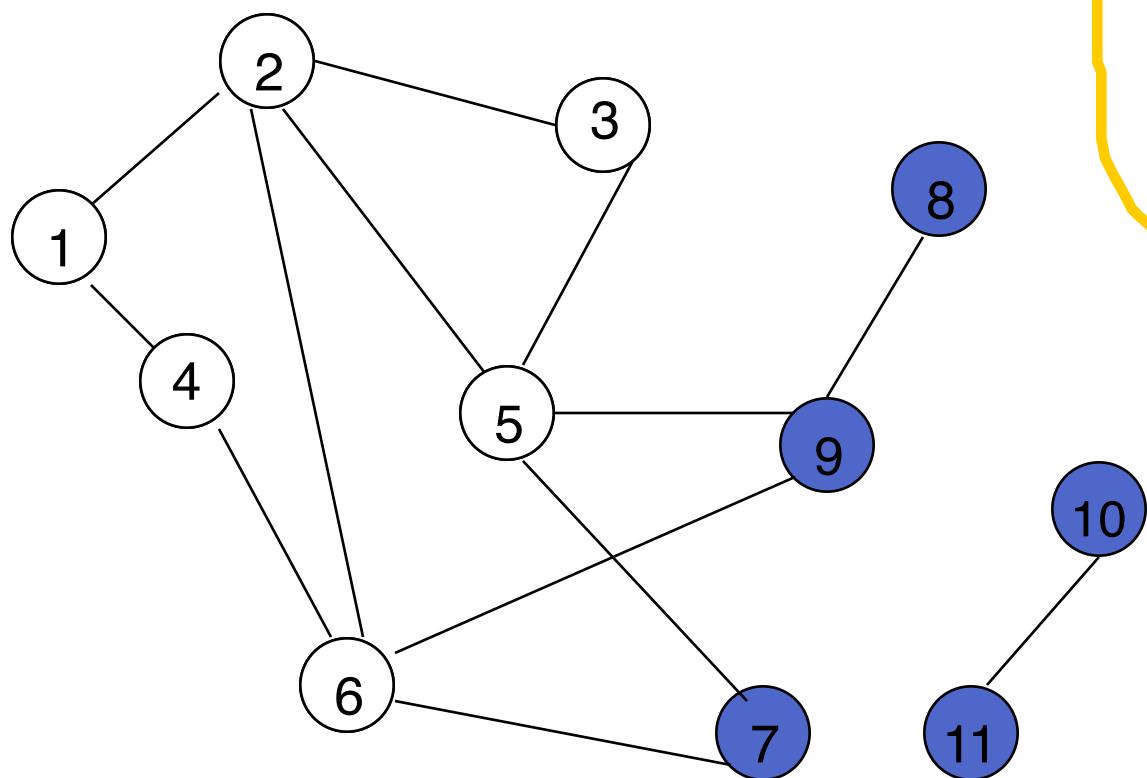
FIFO Queue

5 3 6

Remove 4 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



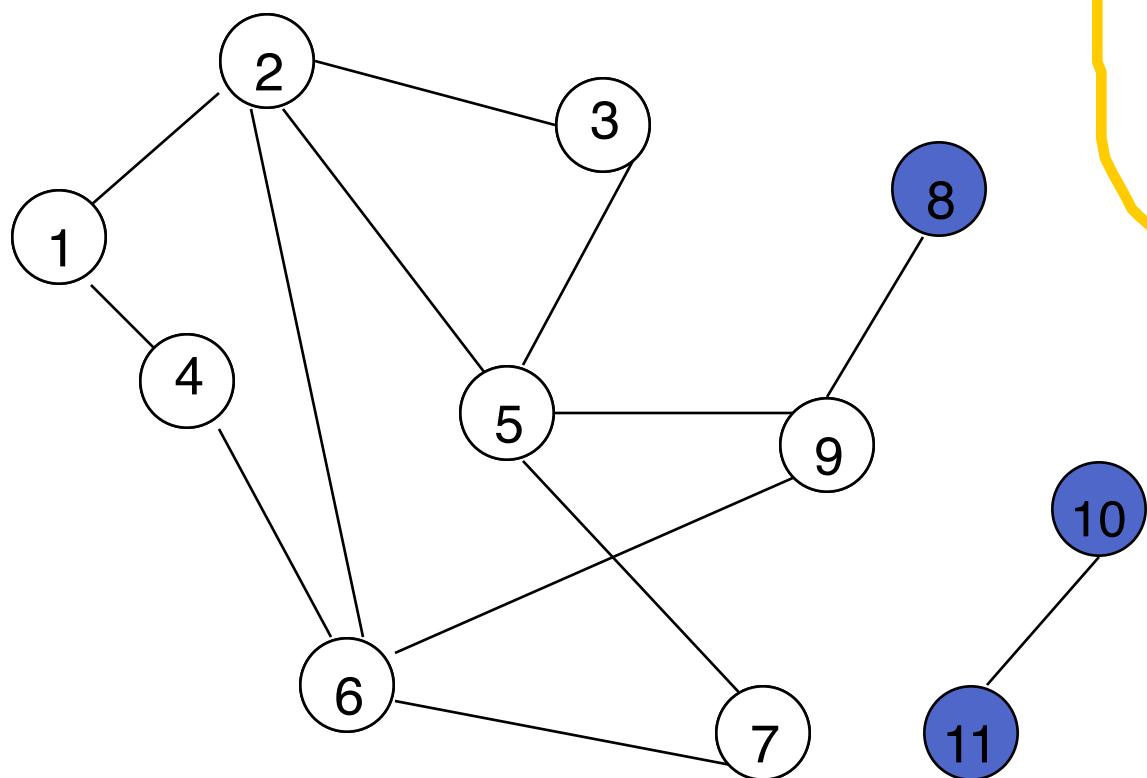
FIFO Queue

5 3 6

Remove 5 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



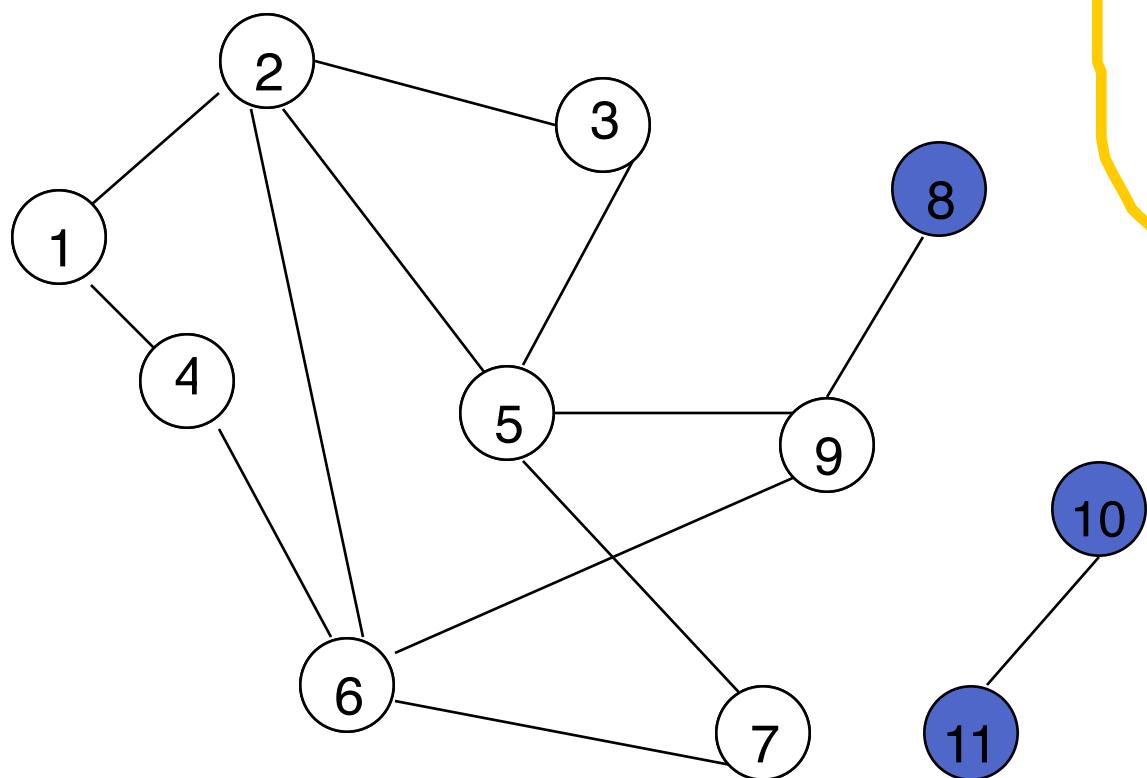
FIFO Queue

3 6 9 7

Remove 5 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



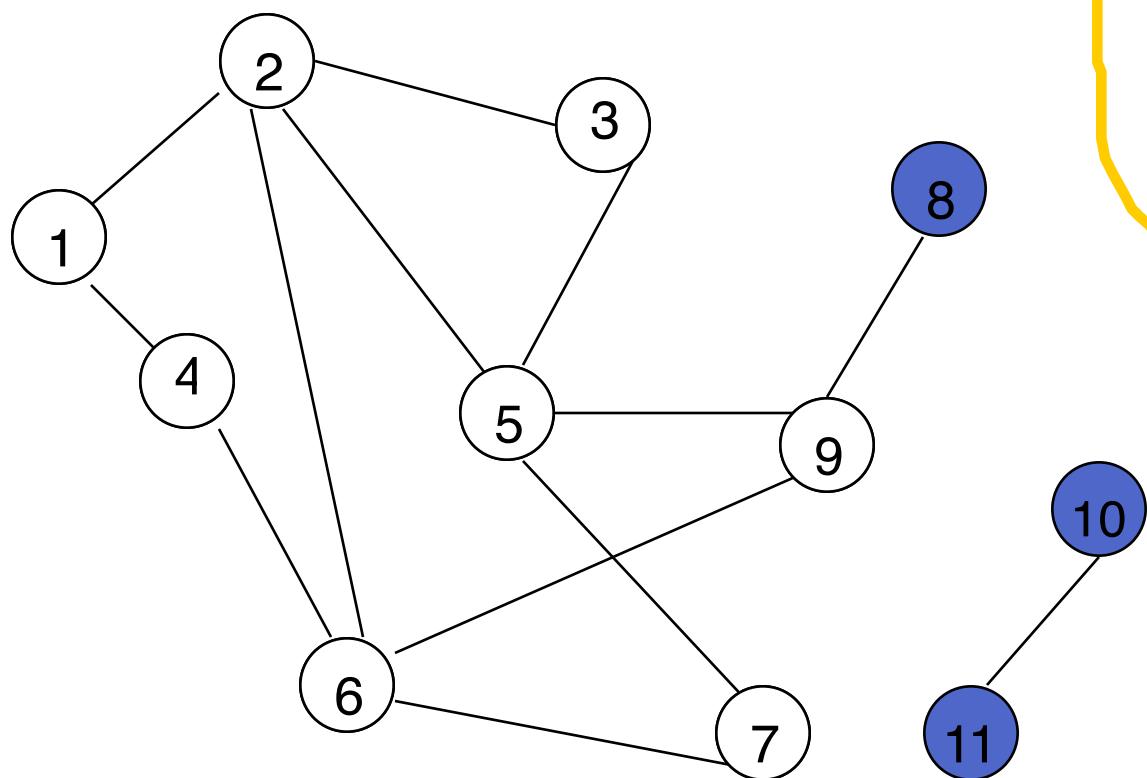
FIFO Queue

3 6 9 7

Remove 3 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



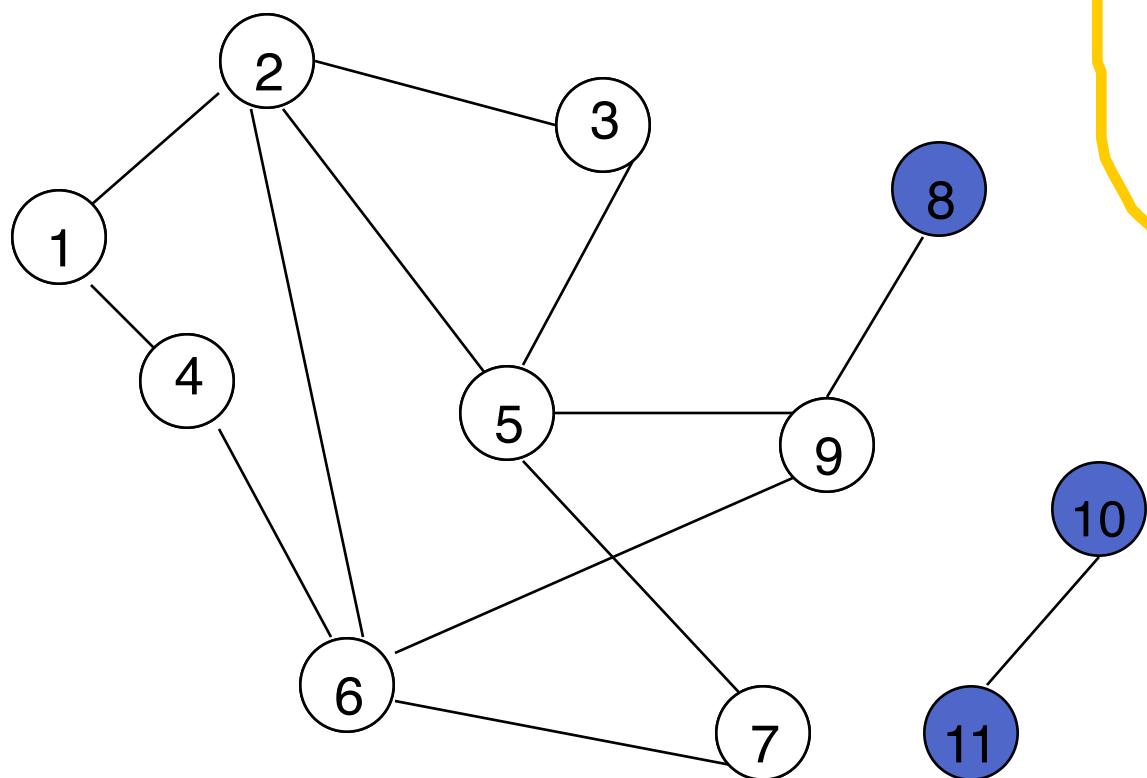
FIFO Queue

6 9 7

Remove 3 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



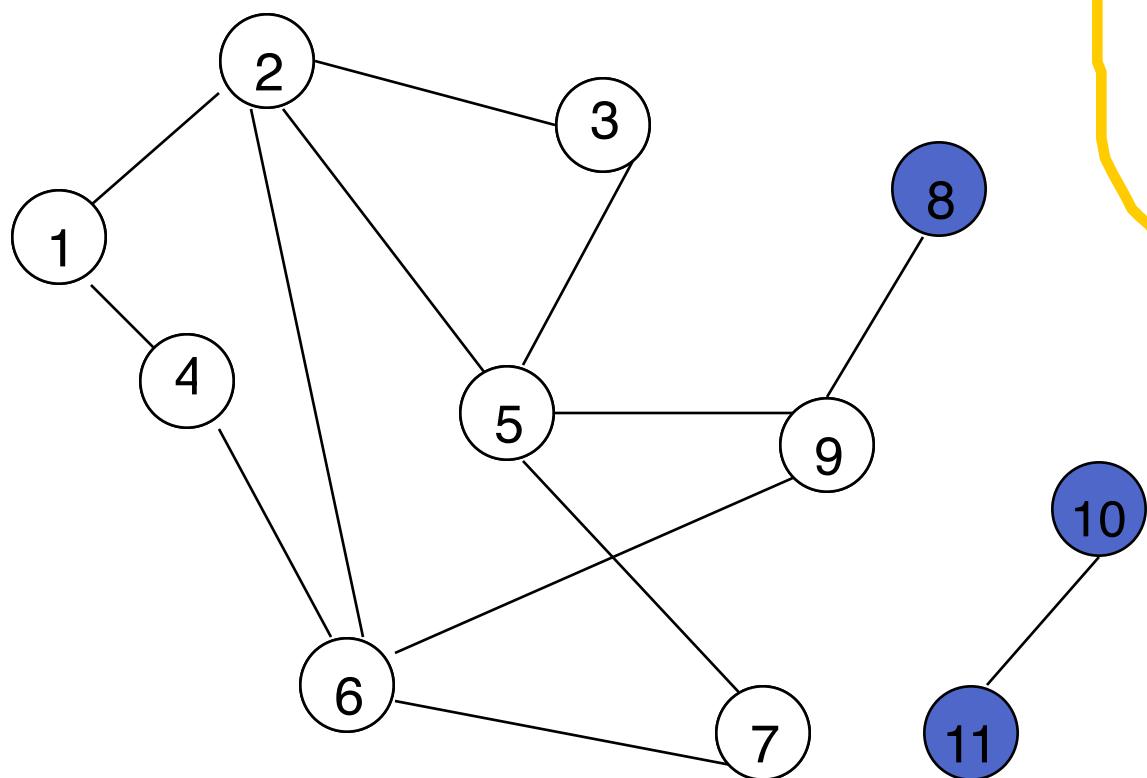
FIFO Queue

6 9 7

Remove 6 from Q; visit adjacent unvisited vertices;
put in Q.



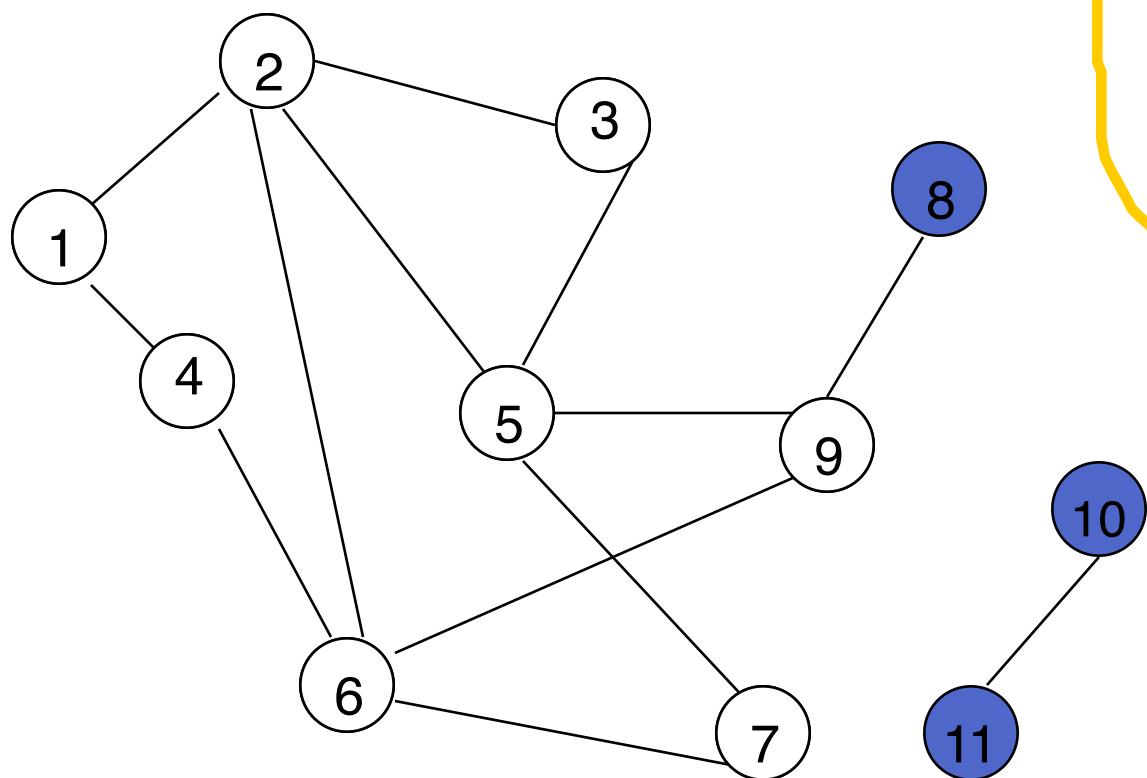
Breadth-First Search Example



Remove 6 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



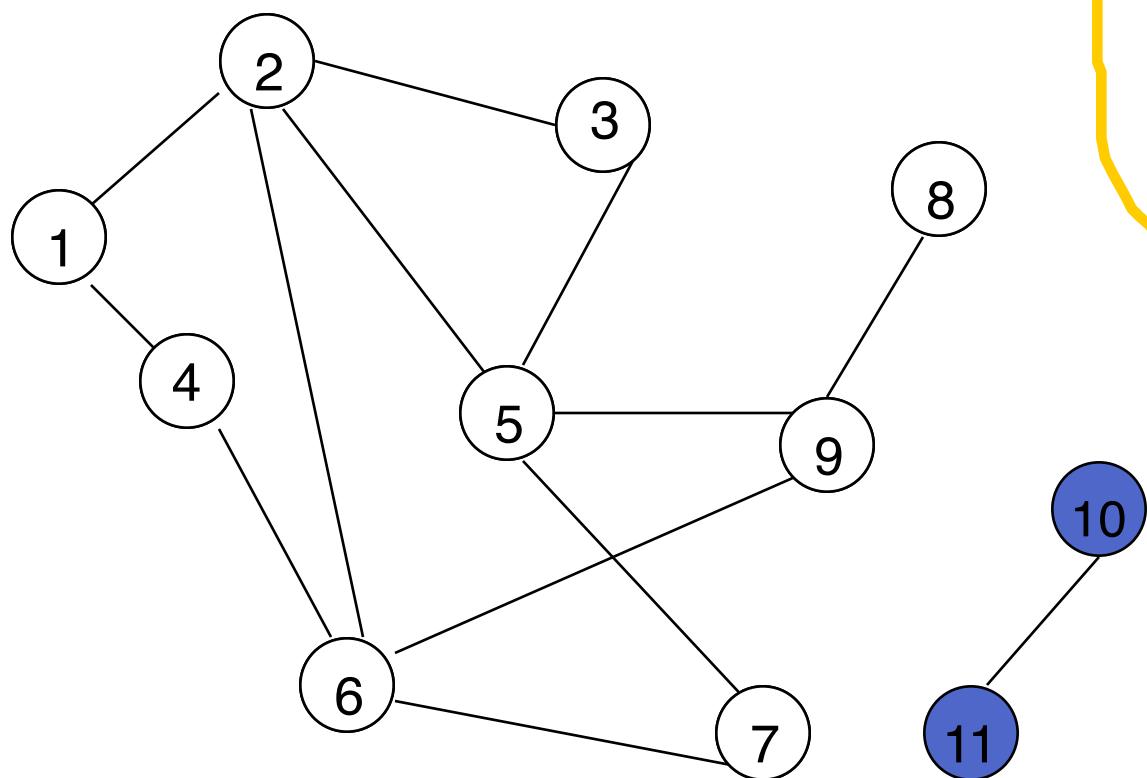
FIFO Queue

9 7

Remove 9 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



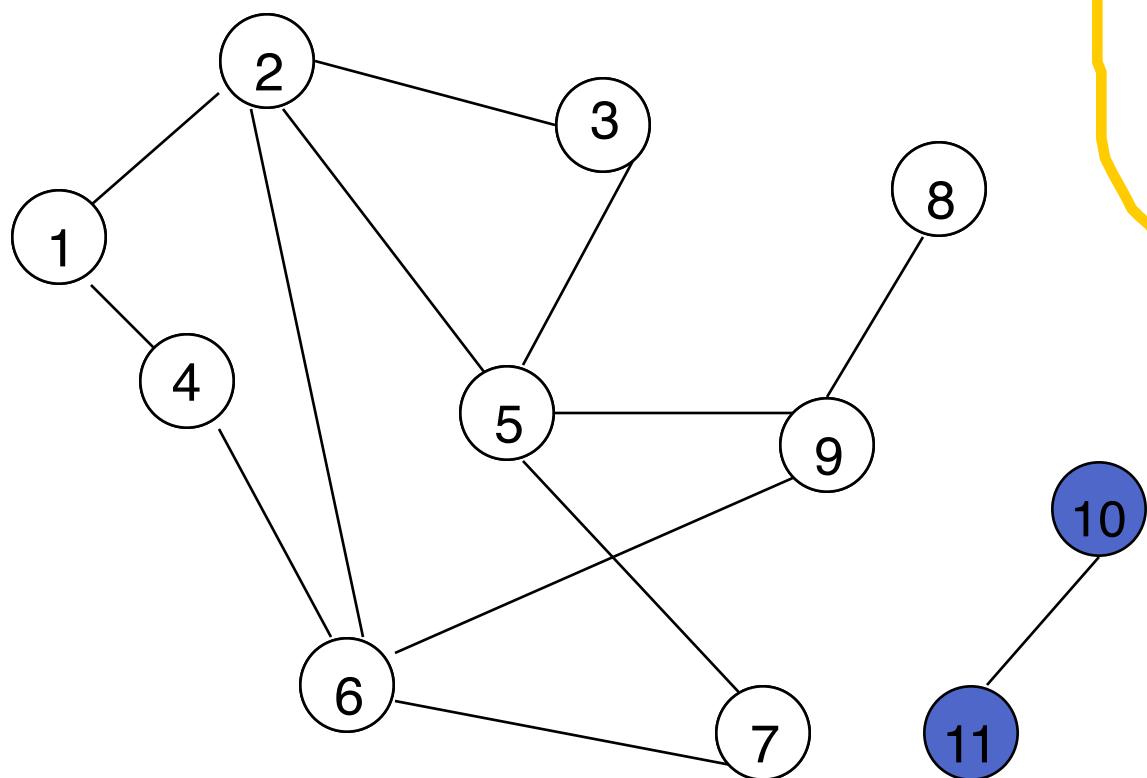
FIFO Queue

7 8

Remove 9 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



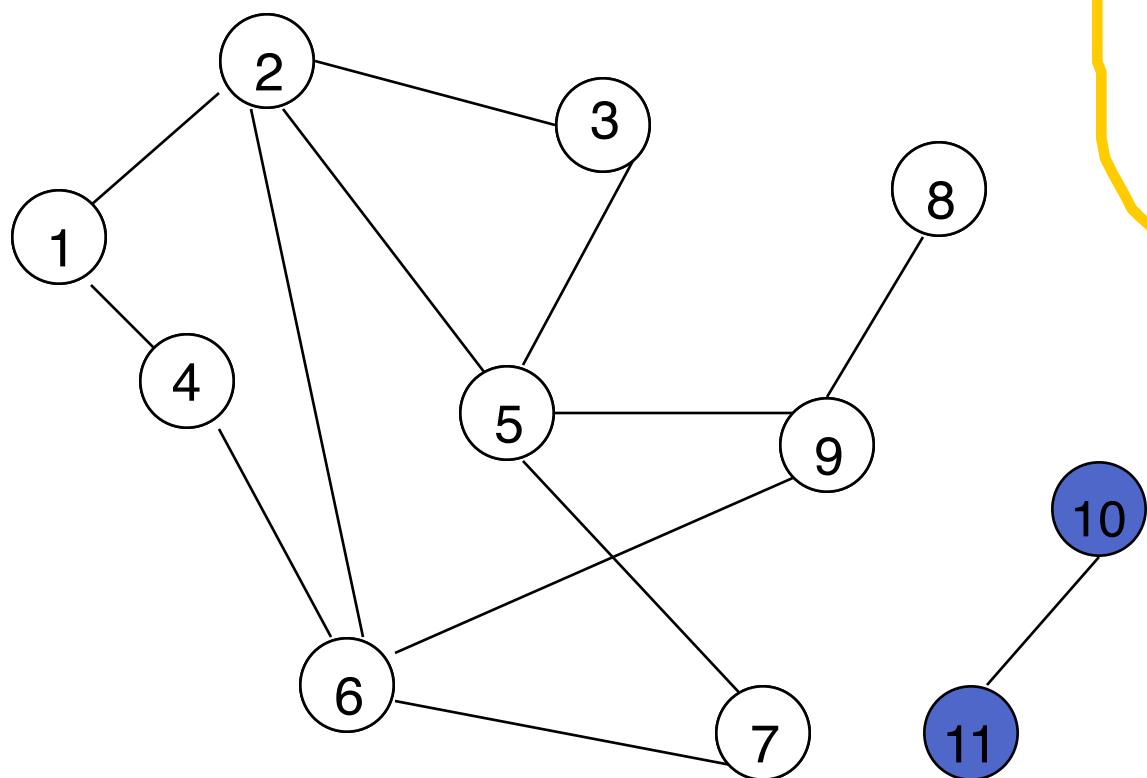
FIFO Queue

7 8

Remove 7 from Q; visit adjacent unvisited vertices;
put in Q.



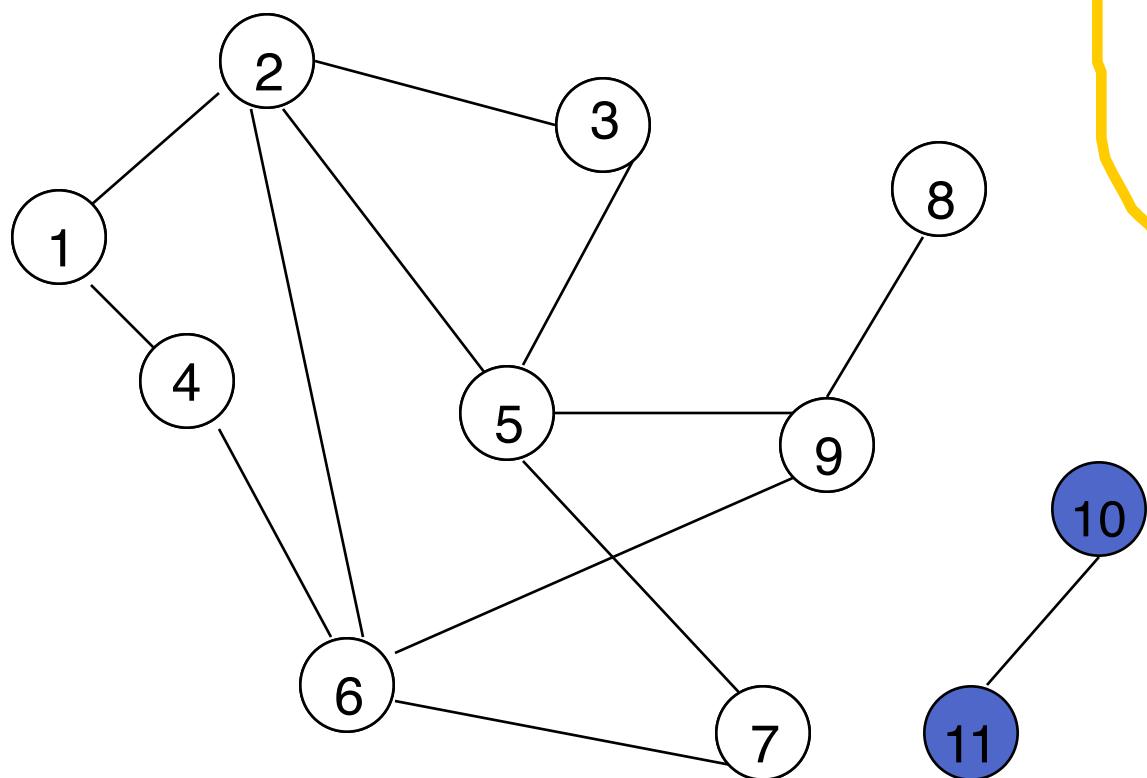
Breadth-First Search Example



Remove 7 from Q; visit adjacent unvisited vertices;
put in Q.



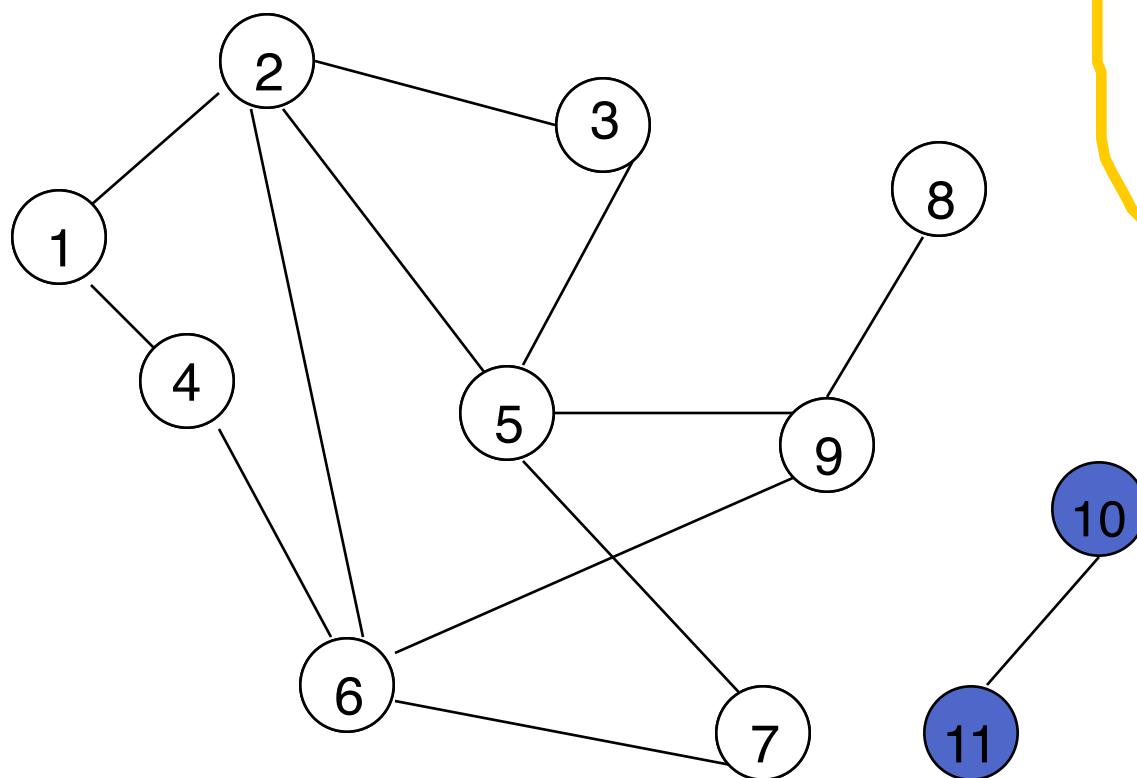
Breadth-First Search Example



Remove 8 from Q; visit adjacent unvisited vertices;
put in Q.



Breadth-First Search Example



Queue is empty. Search terminates.



Breadth-First Search Property

- All vertices reachable from the start vertex (including the start vertex) are visited.



Time Complexity

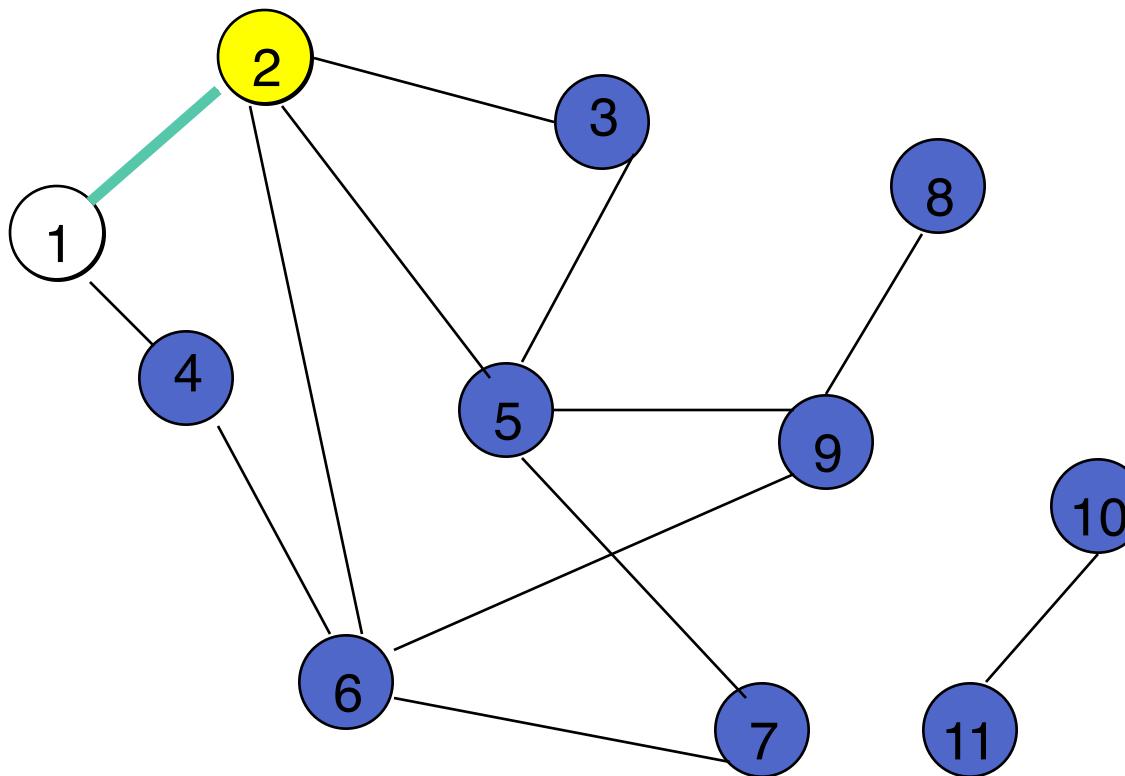
- Each visited vertex is added to (and so removed from) the queue exactly once
- When a vertex is removed from the queue, we examine its adjacent vertices
 - $O(v)$ if adjacency matrix is used, where v is number of vertices in whole graph
 - $O(d)$ if adjacency list is used, where d is *edge degree*
- Total time
 - Adjacency matrix: $O(w.v)$, where w is number of vertices in the *connected component* that is searched
 - Adjacency list: $O(w+f)$, where f is number of edges in the *connected component* that is searched



Depth-First Search

```
depthFirstSearch(v) {  
    Label vertex v as reached;  
    for(each unreached vertex u adjacent to v)  
        depthFirstSearch(u);  
}
```

Depth-First Search

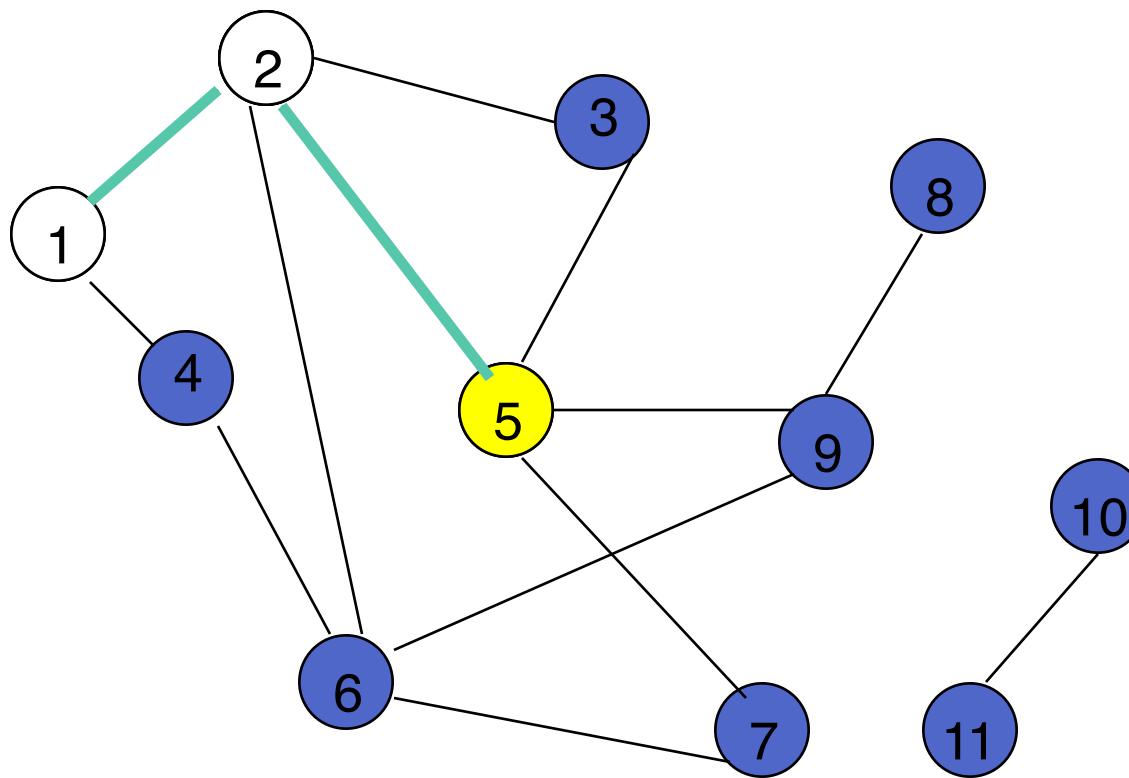


Start search at vertex 1.

Label vertex 1 and do a depth first search from either 2 or 4.

Suppose that vertex 2 is selected.

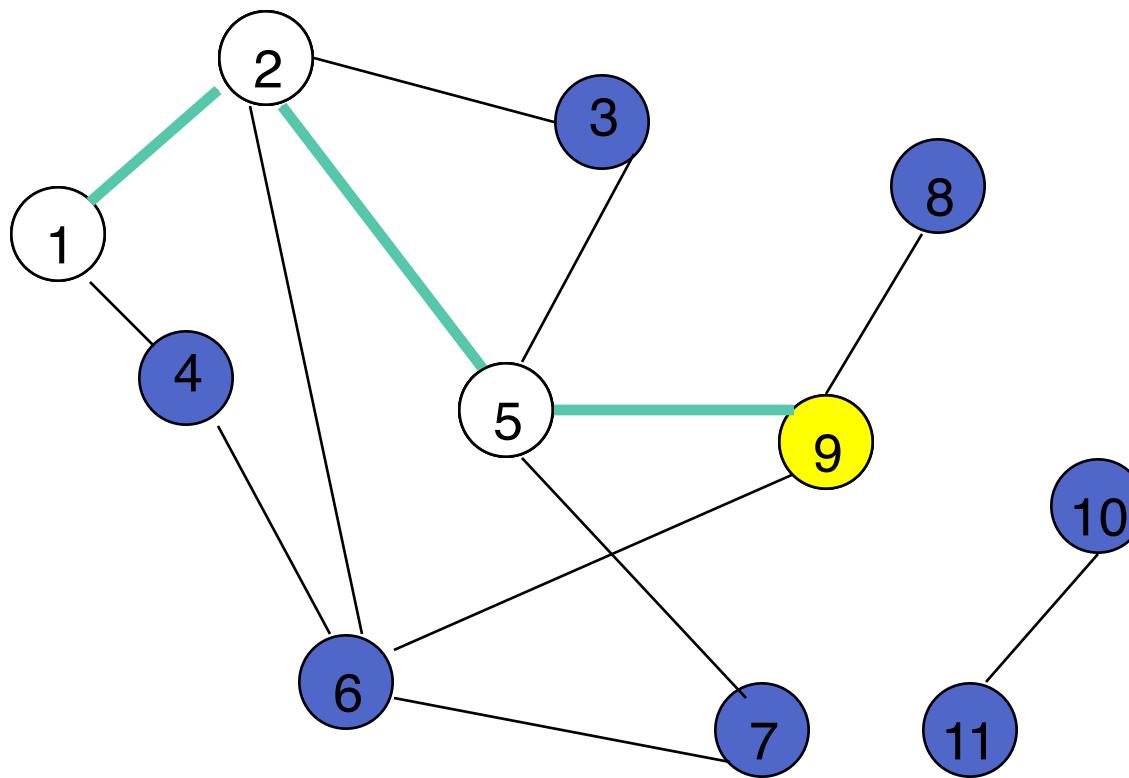
Depth-First Search



Label vertex 2 and do a depth first search from either 3, 5, or 6.

Suppose that vertex 5 is selected.

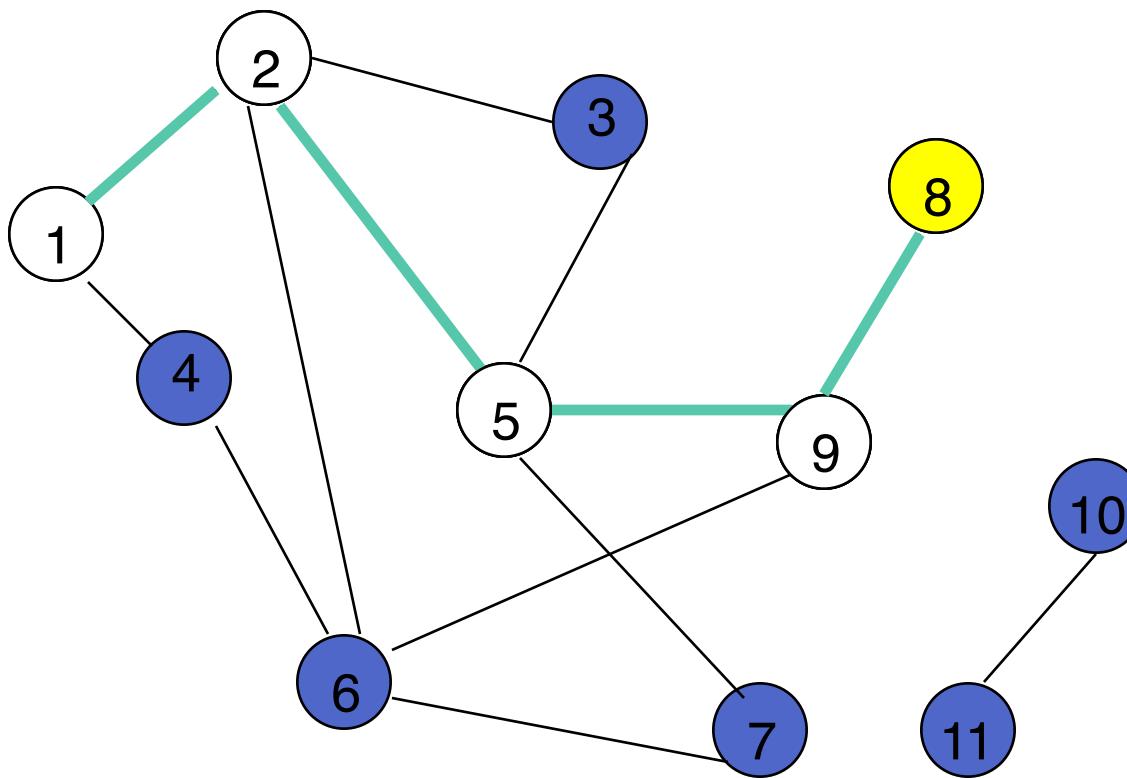
Depth-First Search



Label vertex 5 and do a depth first search from either 3, 7, or 9.

Suppose that vertex 9 is selected.

Depth-First Search

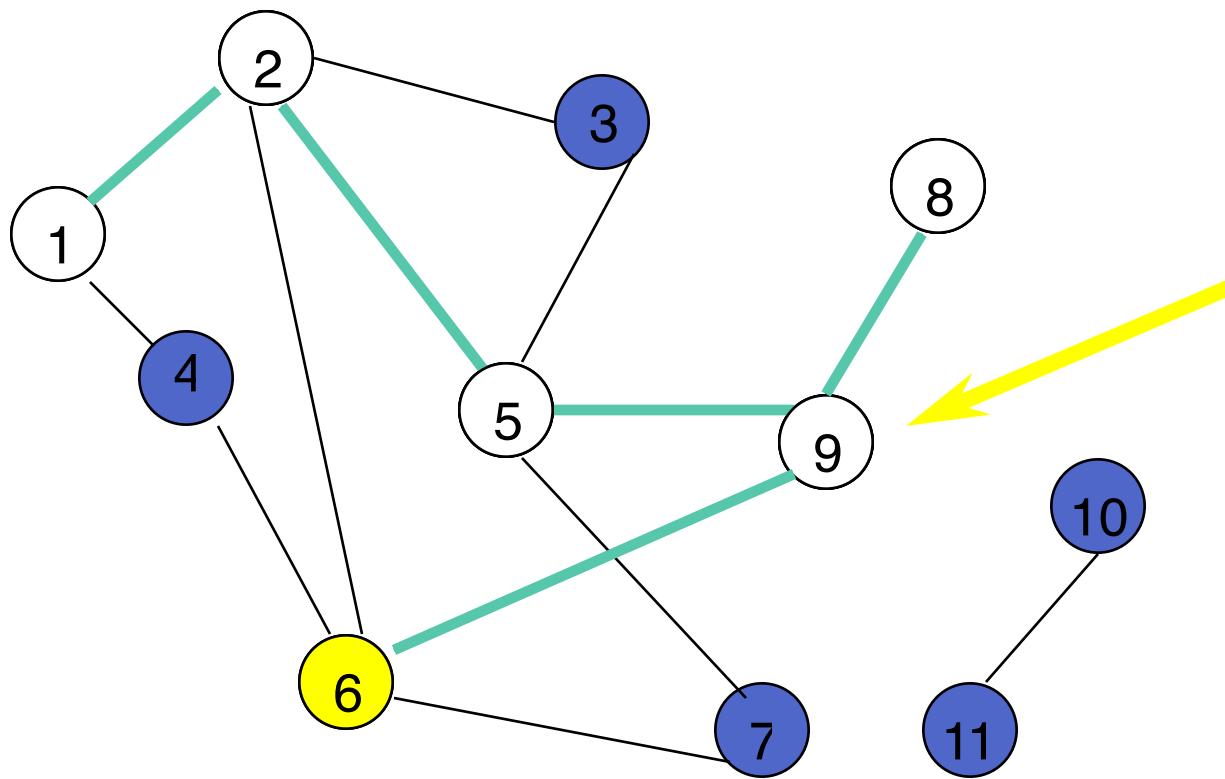


Label vertex 9 and do a depth first search from either 6 or 8.

Suppose that vertex 8 is selected.



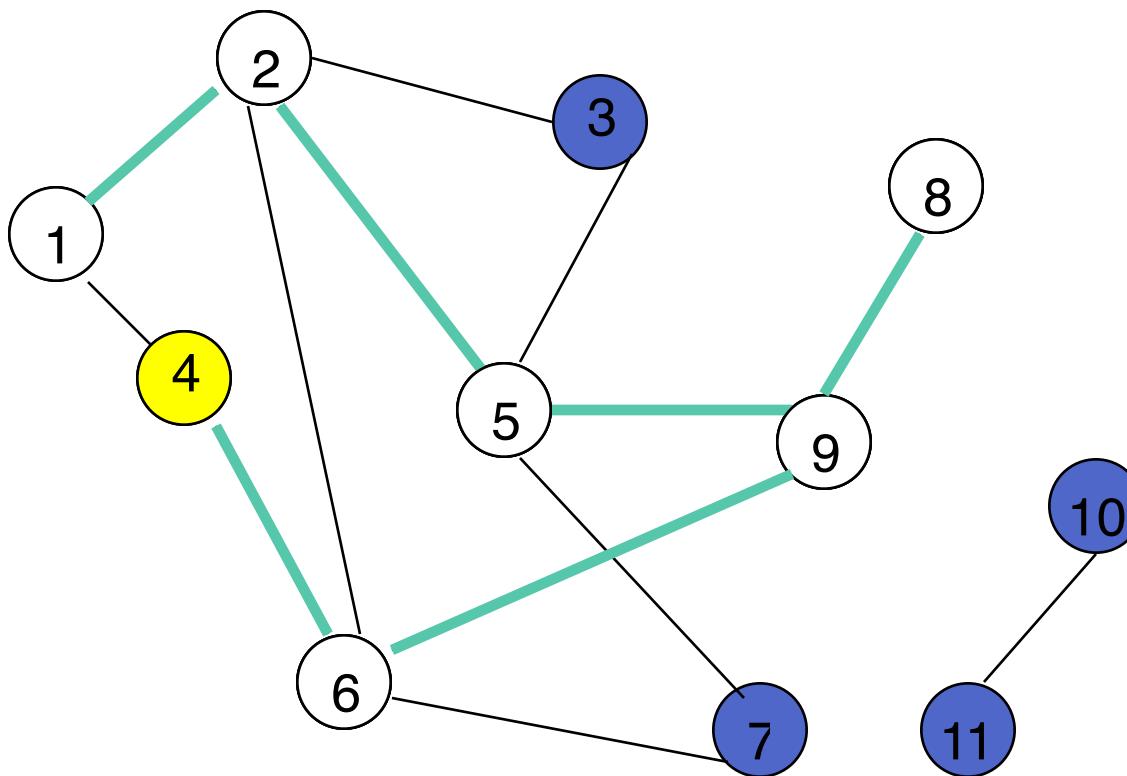
Depth-First Search



Label vertex 8 and return to vertex 9.

From vertex 9 do a $\text{dfs}(6)$

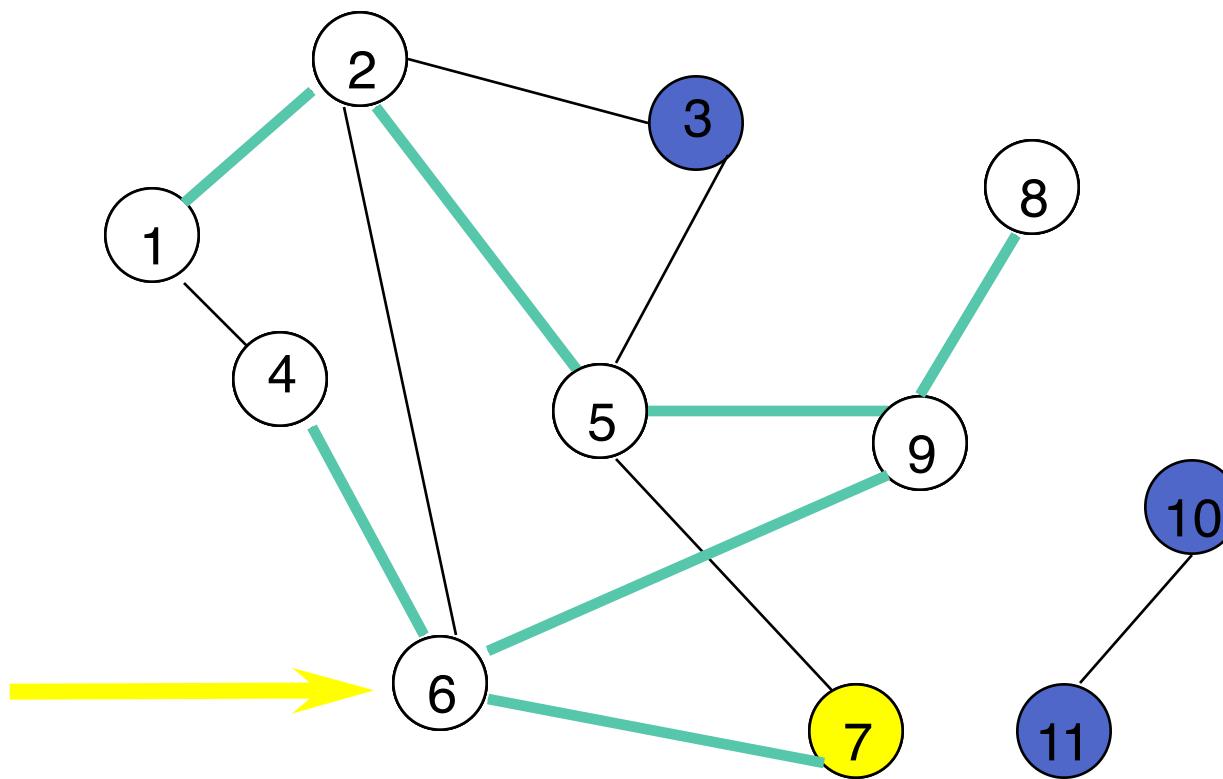
Depth-First Search



Label vertex 6 and do a depth first search from either 4 or 7.

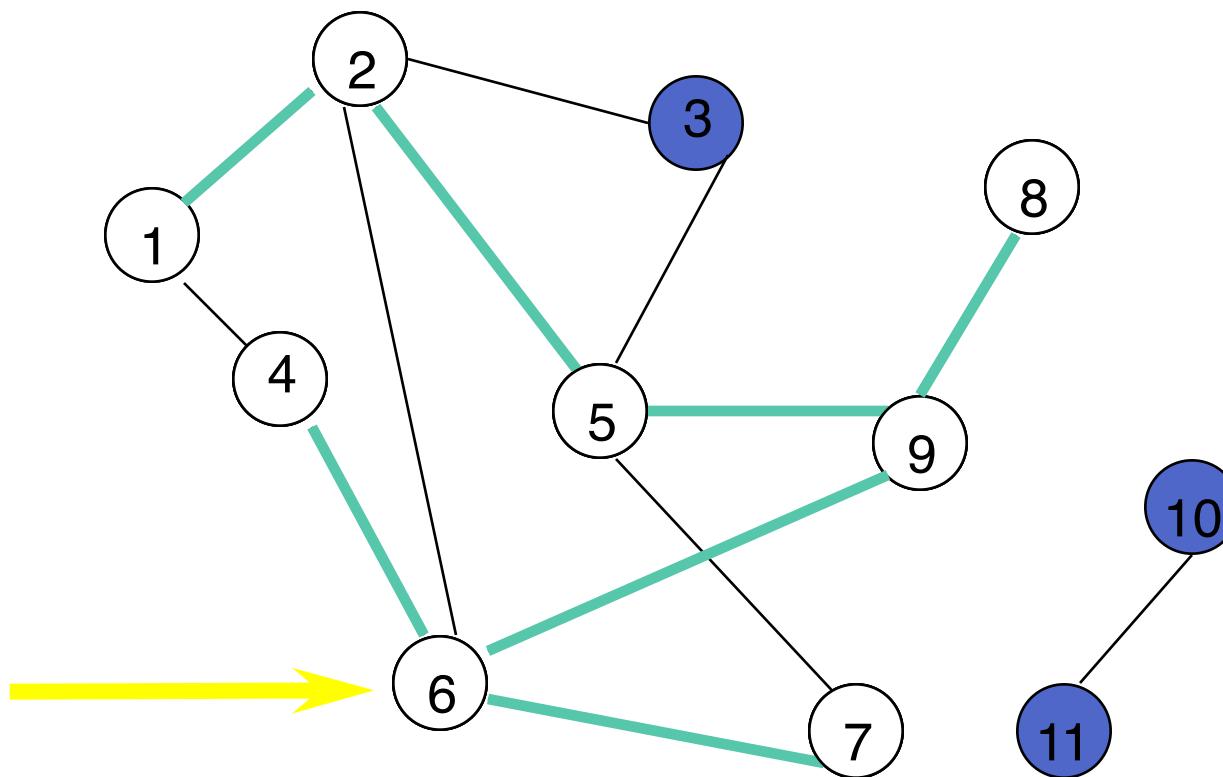
Suppose that vertex 4 is selected.

Depth-First Search



Label vertex 4 and return to 6.
From vertex 6 do a $\text{dfs}(7)$.

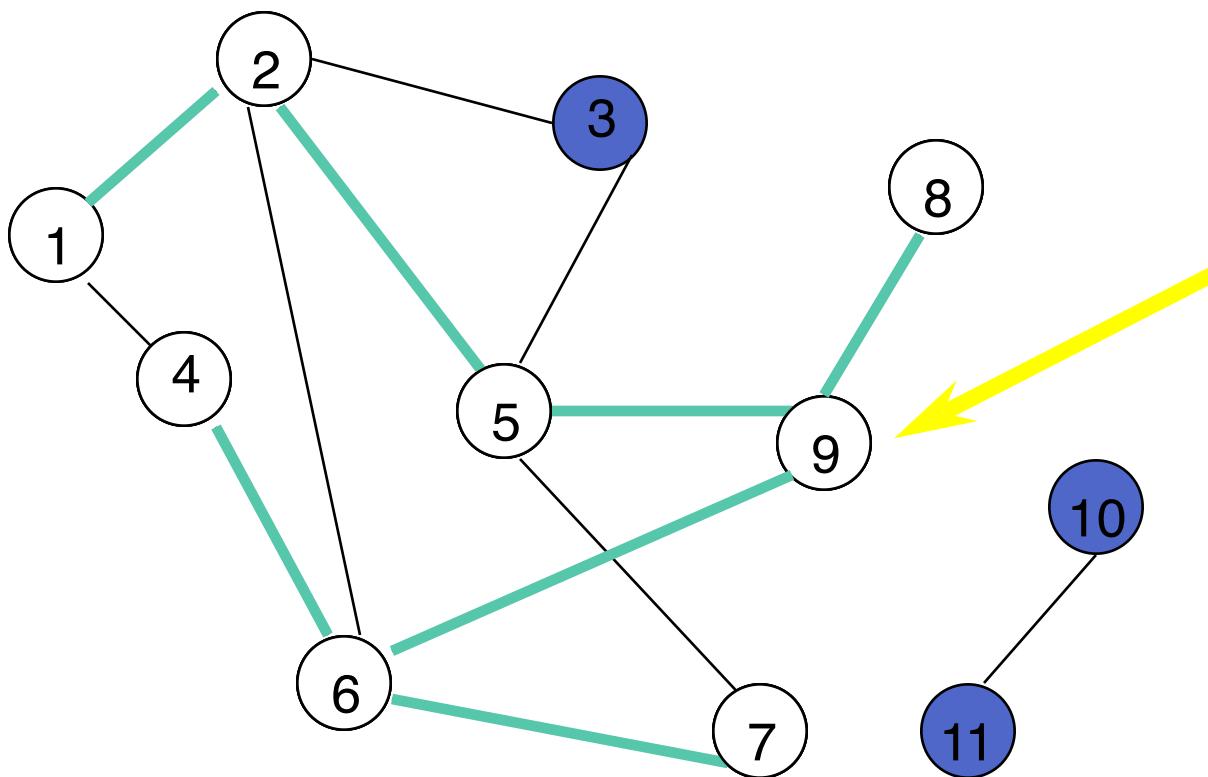
Depth-First Search



Label vertex 7 and return to 6.
Return to 9.



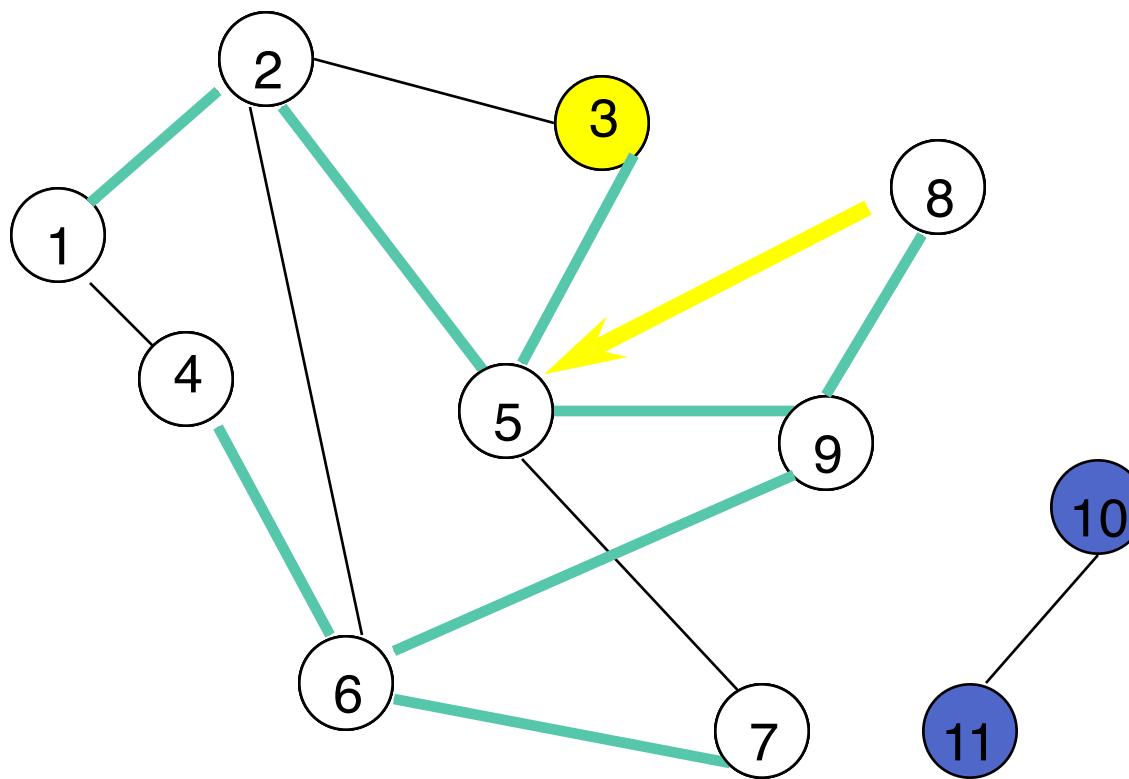
Depth-First Search



Return to 5.

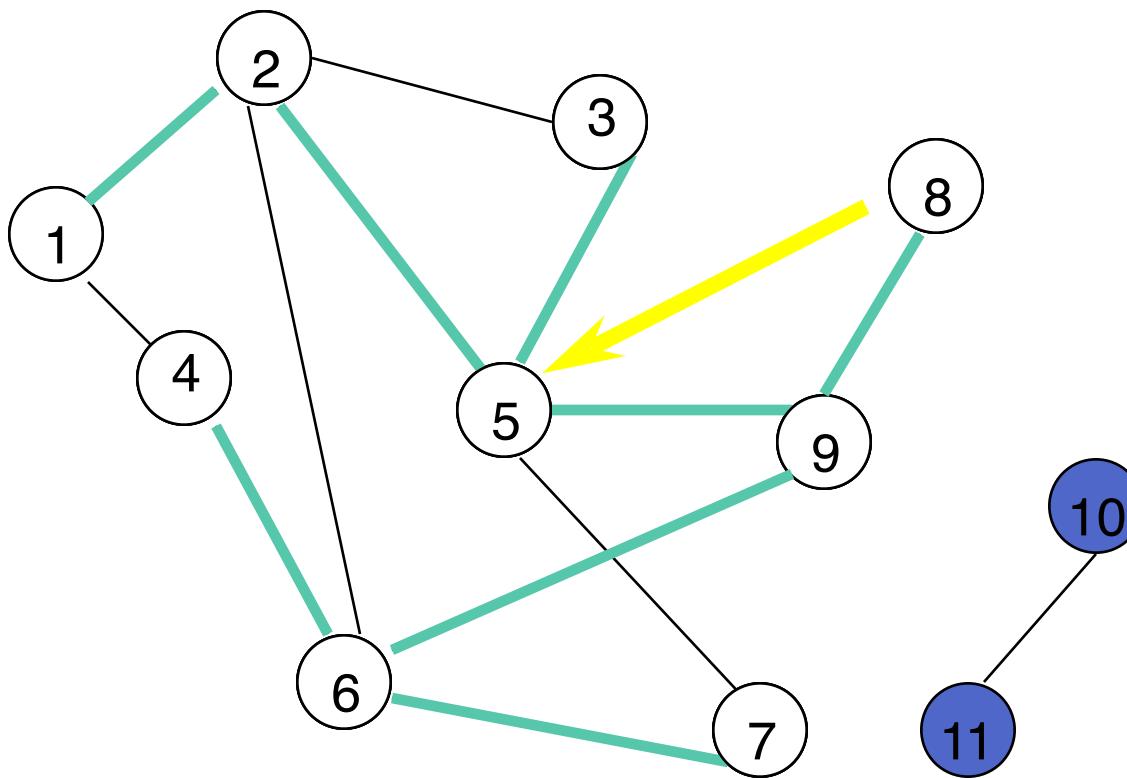


Depth-First Search



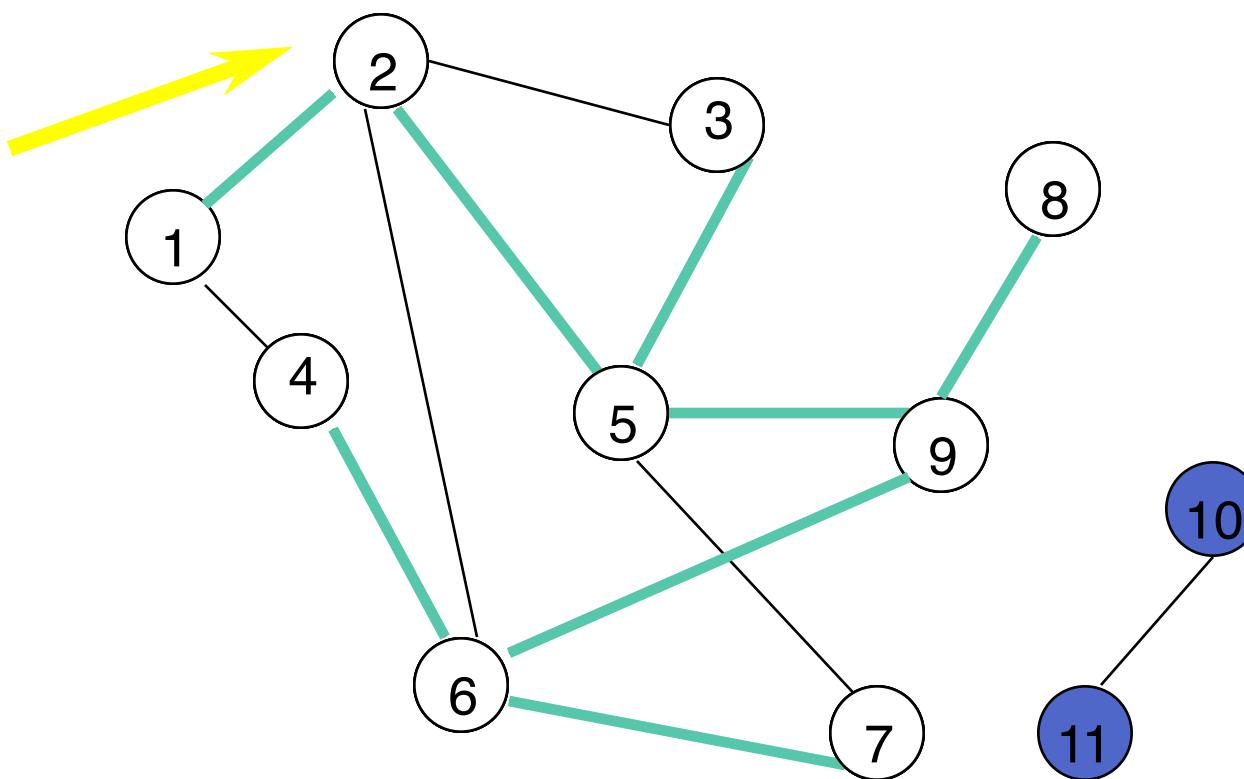
Do a $\text{dfs}(3)$.

Depth-First Search



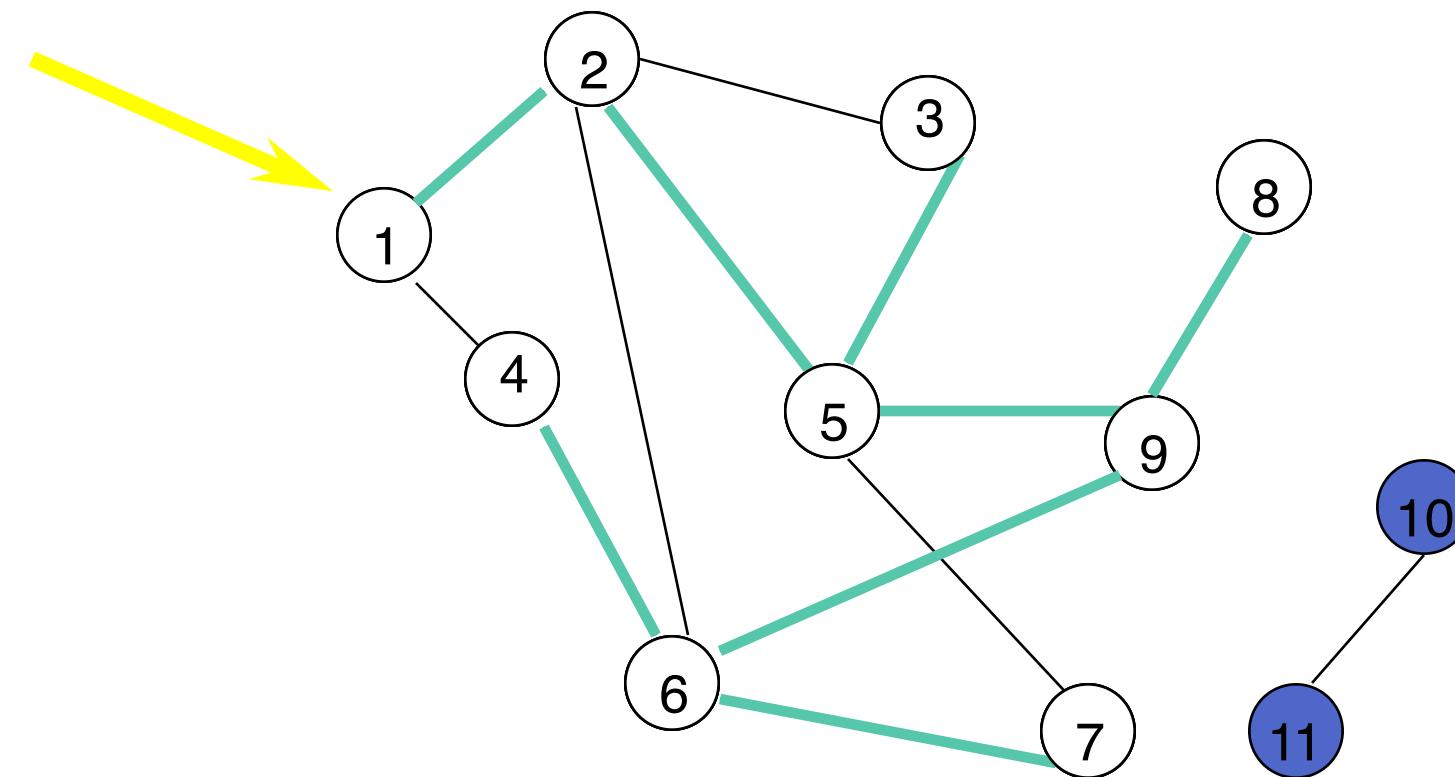
Label 3 and return to 5.
Return to 2.

Depth-First Search



Return to 1.

Depth-First Search



Return to invoking method.

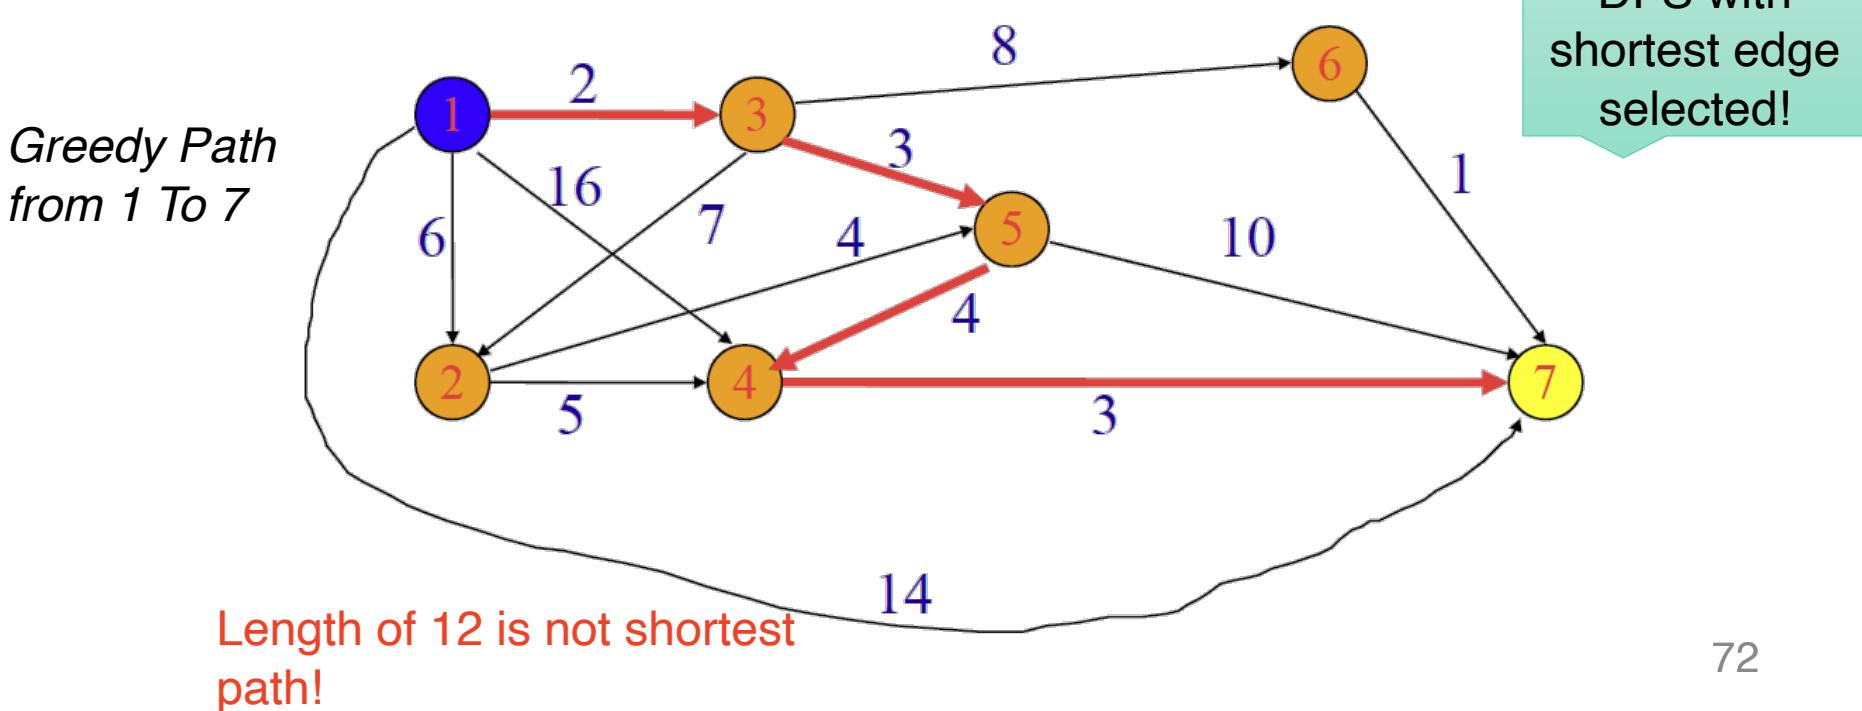


DFS Properties

- DFS has same time complexity as BFS
- DFS requires $O(v)$ memory for recursive function stack while BFS requires $O(v)$ queue capacity, $v=\# \text{vertices}$
- Same properties with respect to path finding, connected components, and spanning trees.
 - Edges used to reach unlabelled vertices define a depth-first spanning tree when the graph is connected.
- One is better than the other for some problems, e.g.
 - When searching, if the item is far from source (leaves), then DFS may locate it first, and vice versa for BFS
 - BFS traverses vertices at same distance (level) from source
 - DFS can be used to detect cycles (revisits of vertices in current stack)

Shortest Path: Single source, single destination

- Possible greedy algorithm
 - Leave source vertex using *shortest edge*
 - Leave new vertex using *cheapest edge*, to reach an *unvisited vertex*
 - Continue until destination is reached



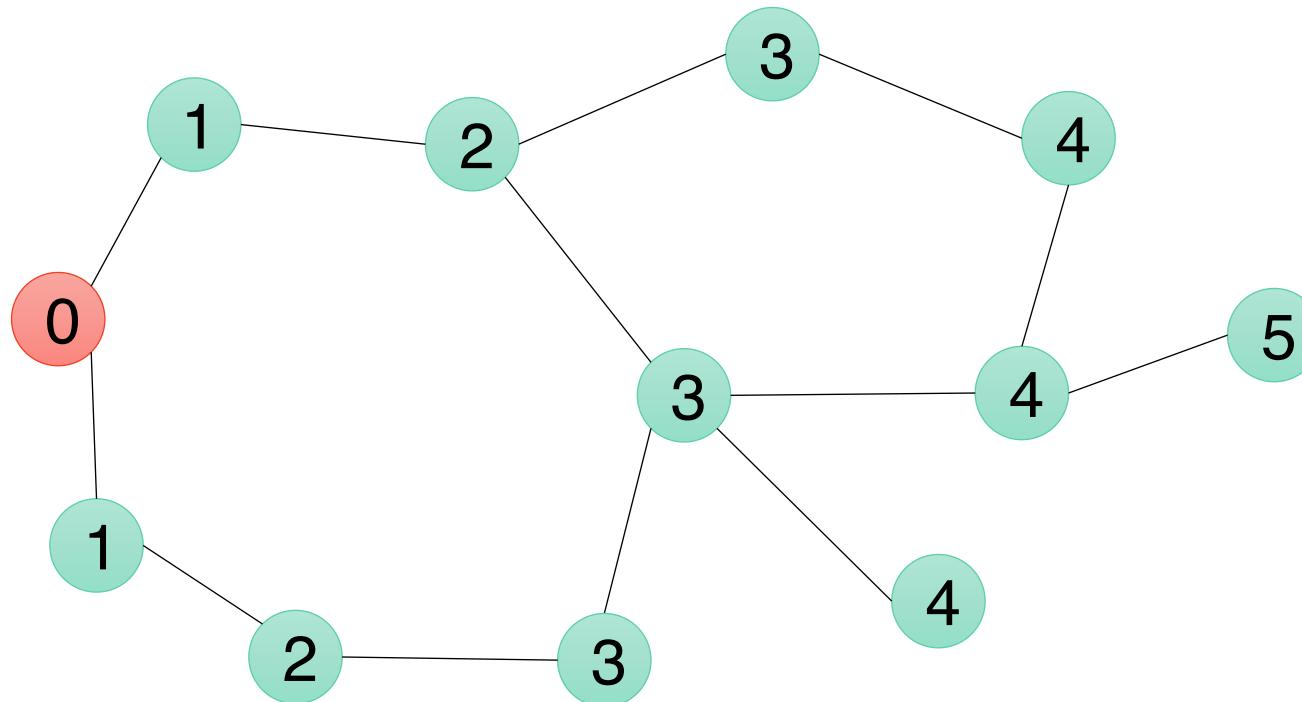


Single Source Shortest Path

- Shortest distance from one source vertex to all destination vertices
- Is there a simple way to solve this?
- ...Say if you had a unit-weighted graph?
- Just do Breadth First Search (BFS)! 😊

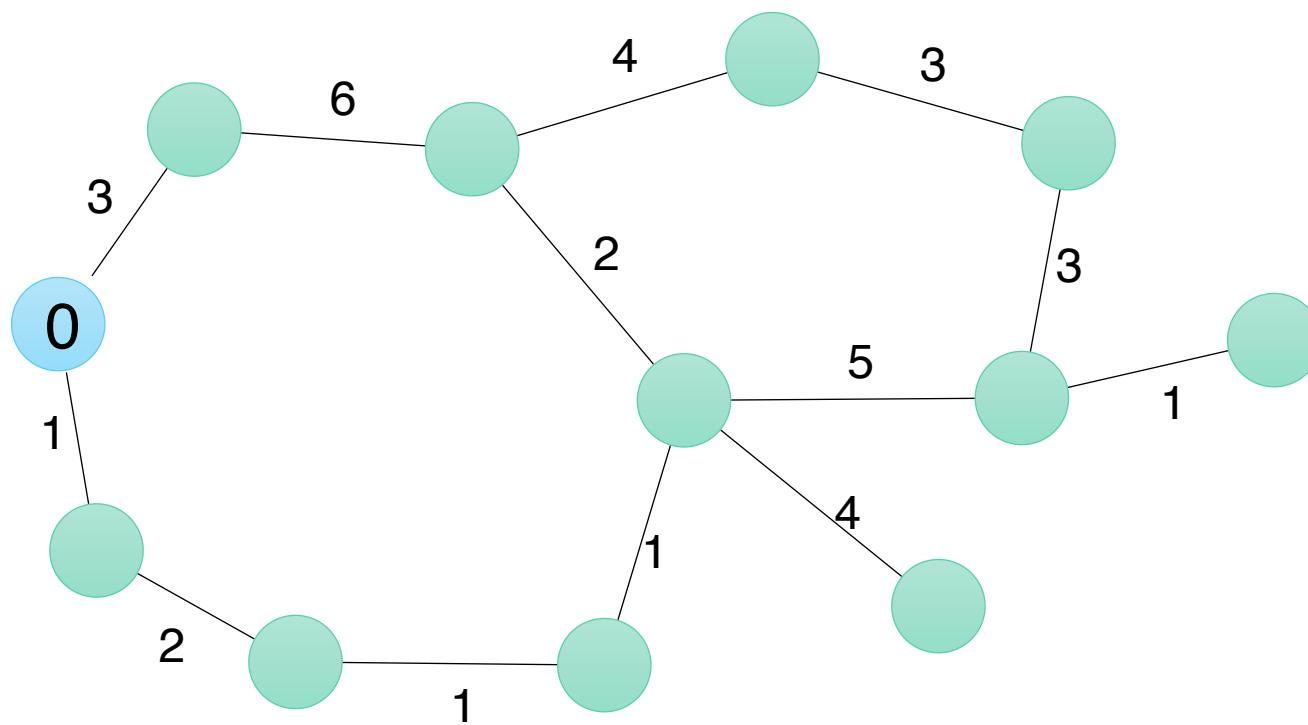


SSSP: BFS on Unweighted Graphs



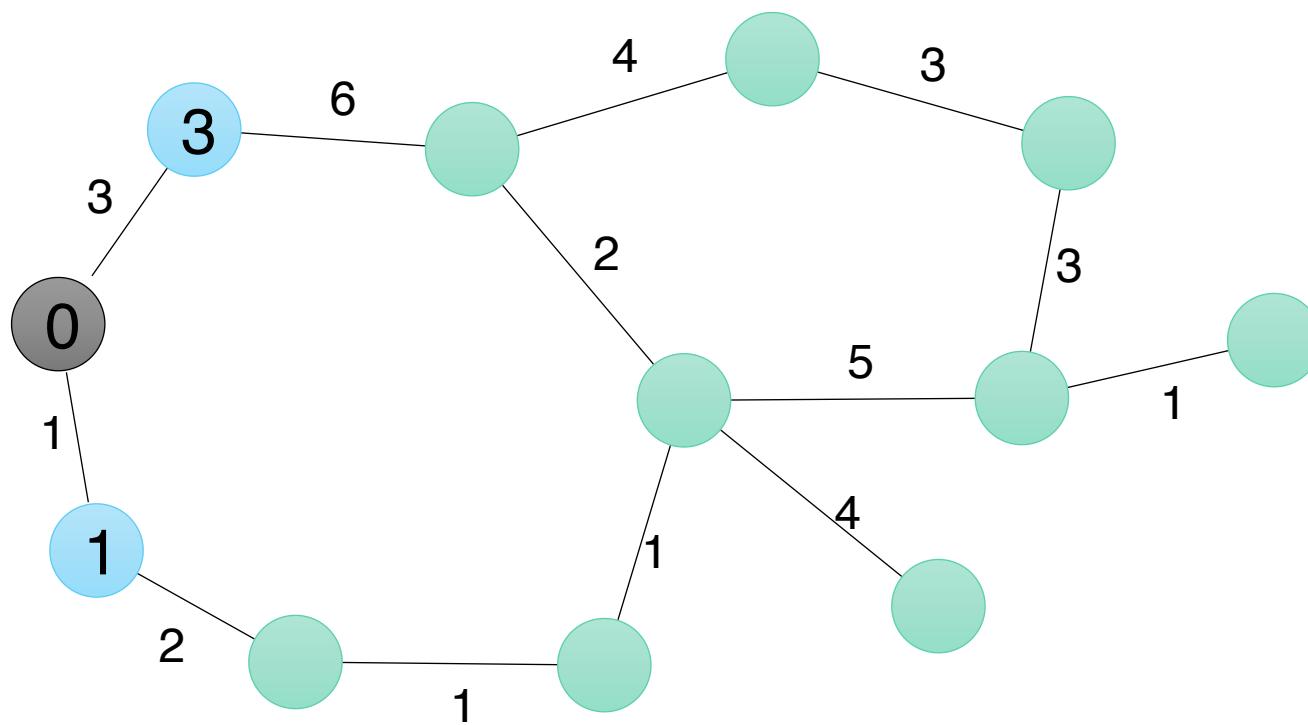


SSSP: BFS on Weighted Graphs?



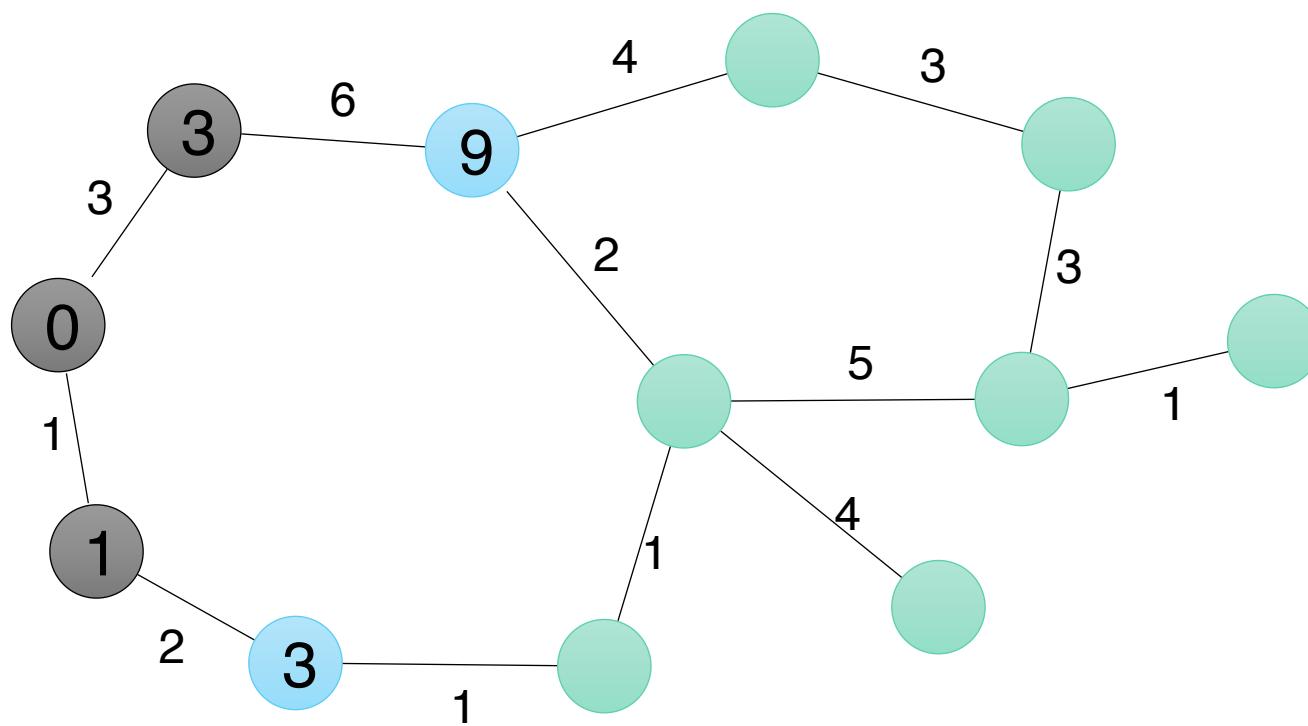


SSSP: BFS on Weighted Graphs?



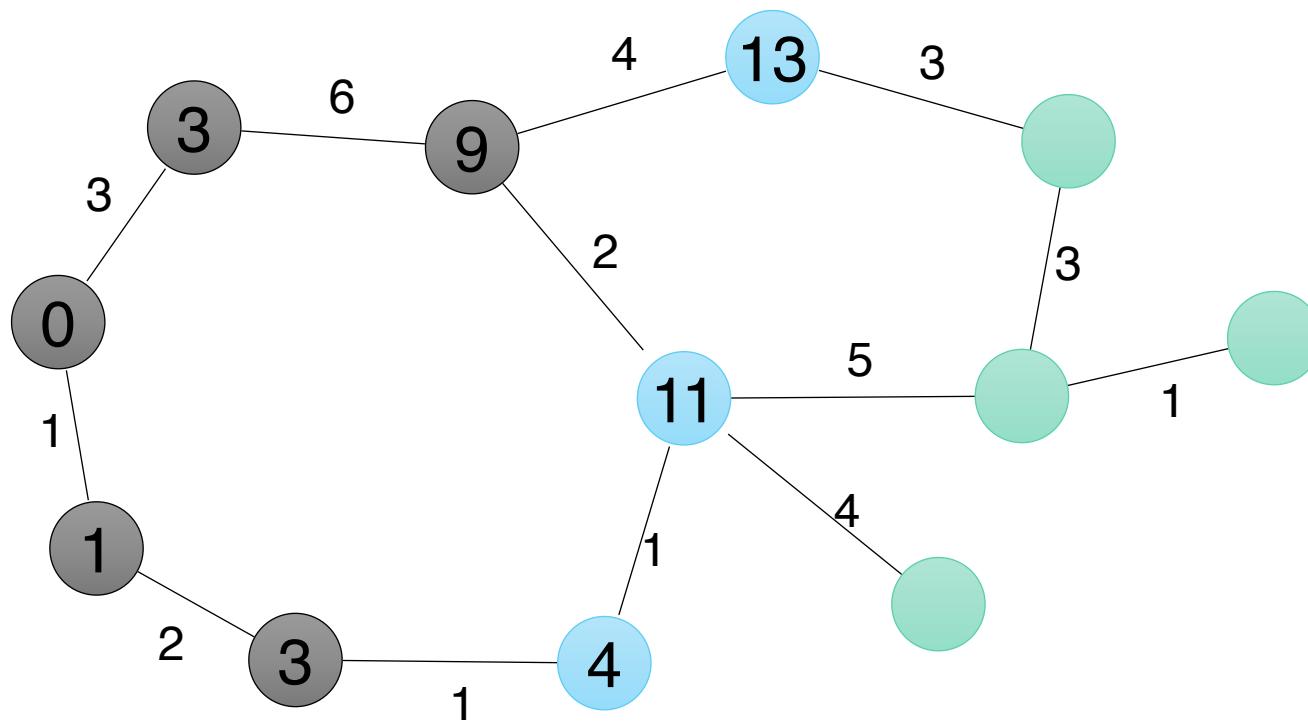


SSSP: BFS on Weighted Graphs?



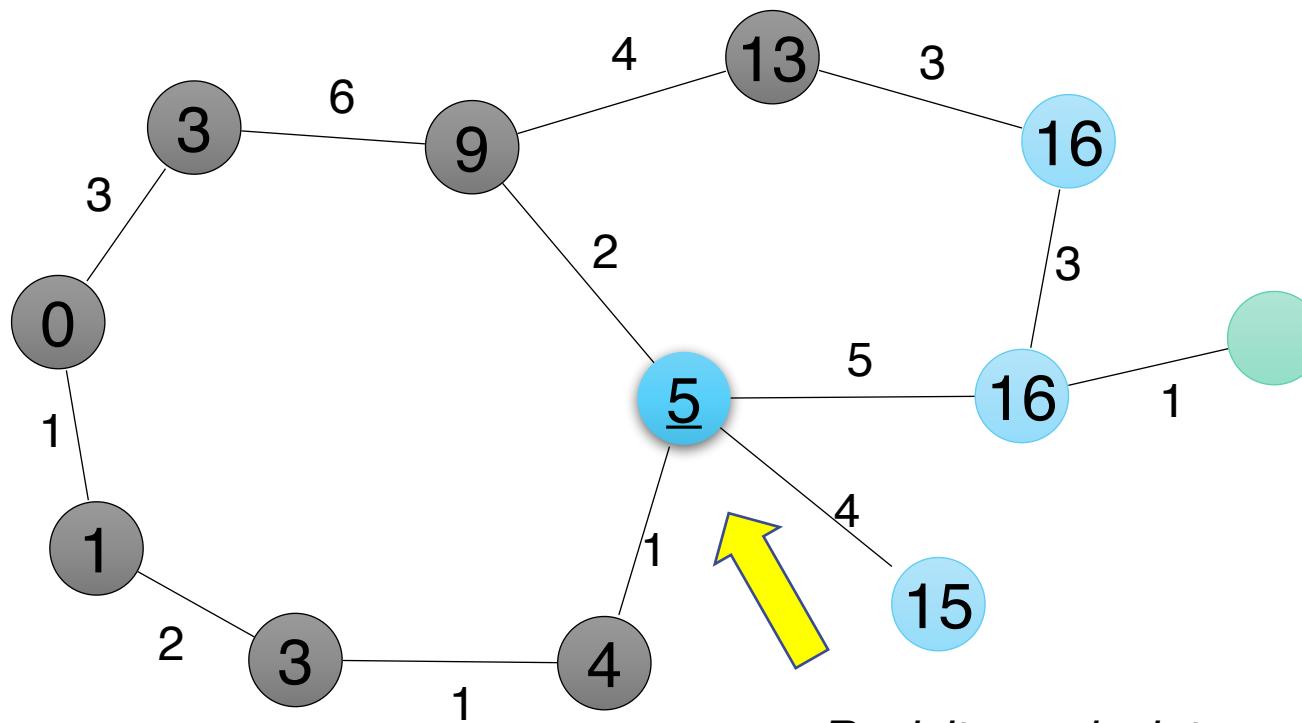


SSSP: BFS on Weighted Graphs?





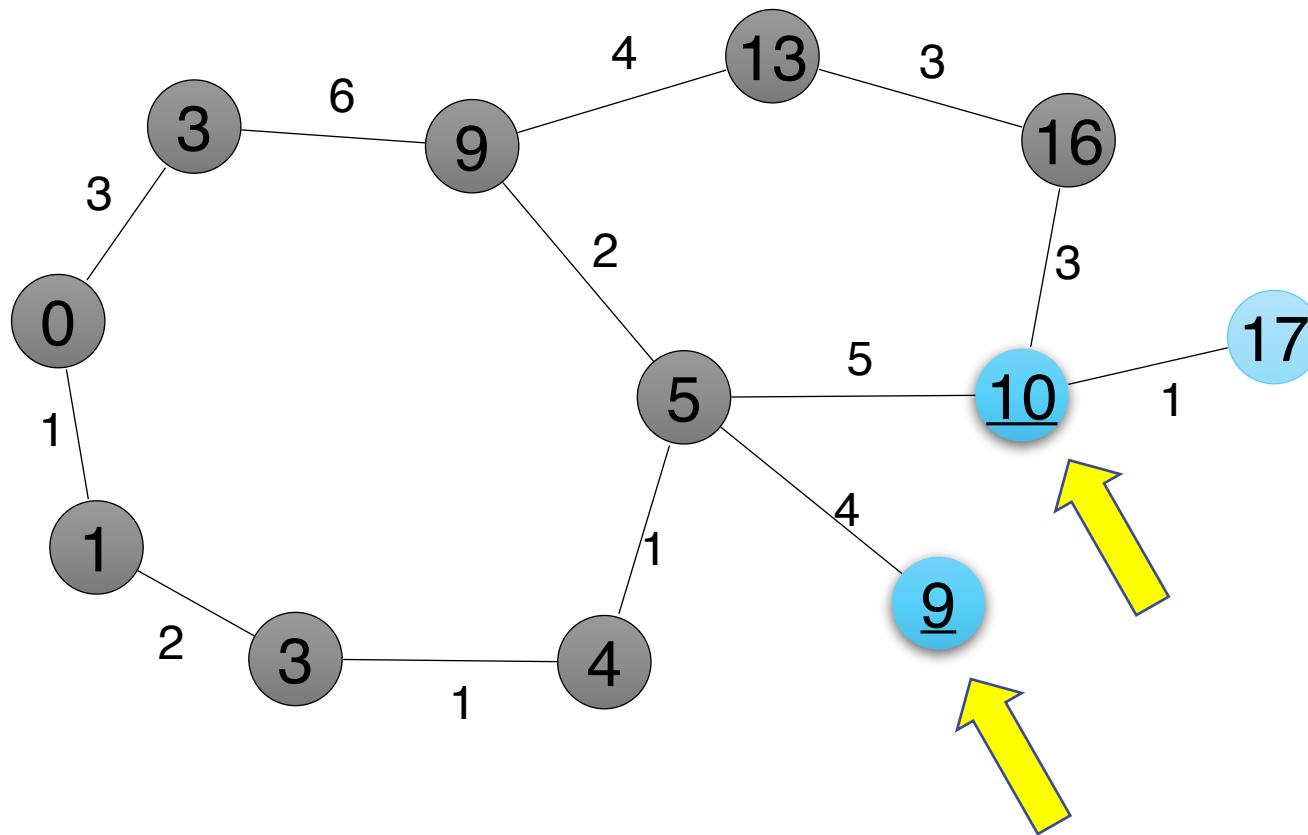
SSSP: BFS on Weighted Graphs?



*Revisit, recalculate, re-propagate...
cascading effect*

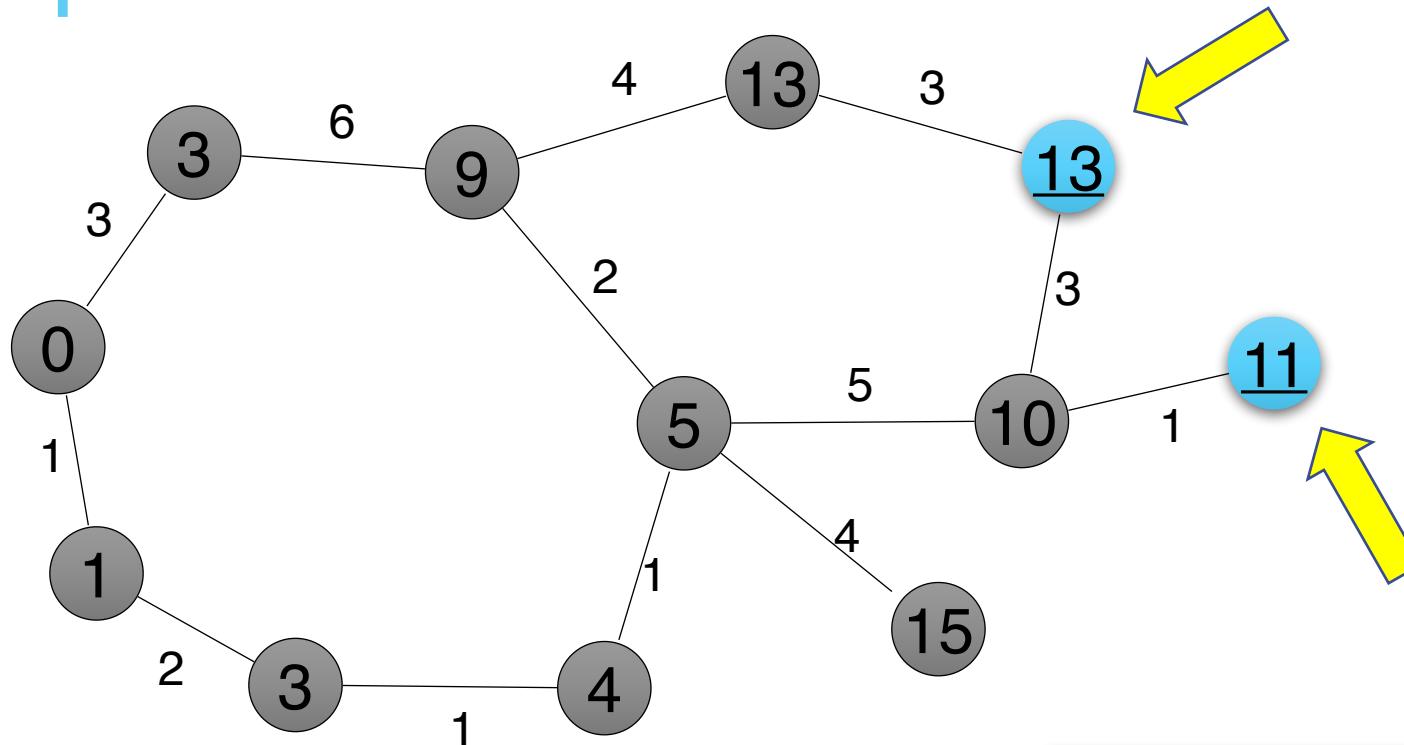


SSSP: BFS on Weighted Graphs?





SSSP: BFS on Weighted Graphs?



BFS with revisits is not efficient. Can we be smart about order of visits?



Dijkstra's Single Source Shortest Path (SSSP)

- Similar to BFS, but prioritise the vertices to visit next
 - Pick “unvisited” vertex with shortest distance from source
- Do not revisit vertices that have already been visited
 - Avoids false propagation of distances

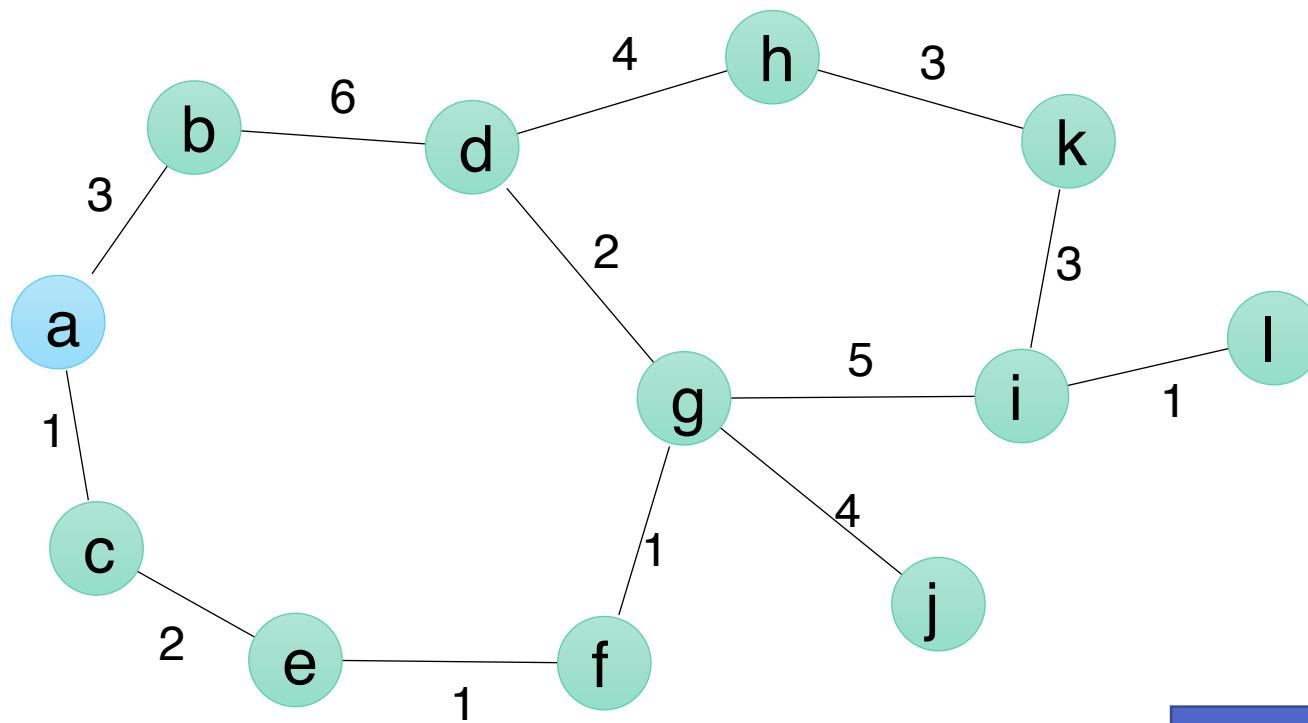


Dijkstra's Single Source Shortest Path (SSSP)

- Let $w[u,v]$ be array with weight of edge from u to v
- Initialise distance vector $d[]$ for all vertices to **infinity**, except for source which is set to **0**
- Add source vertex to container **Q**
- while(**Q** is not empty)
 - Remove **u** from **Q** such that **d[u] is the smallest in Q**
 - Mark **u** as **visited**
 - for each **v** adjacent to **u** that is not **visited**
 - Add **v** to container **Q** if not present in it
 - Set $d[v] = \min(d[v], d[u] + w[u,v])$
- **Runtime?**



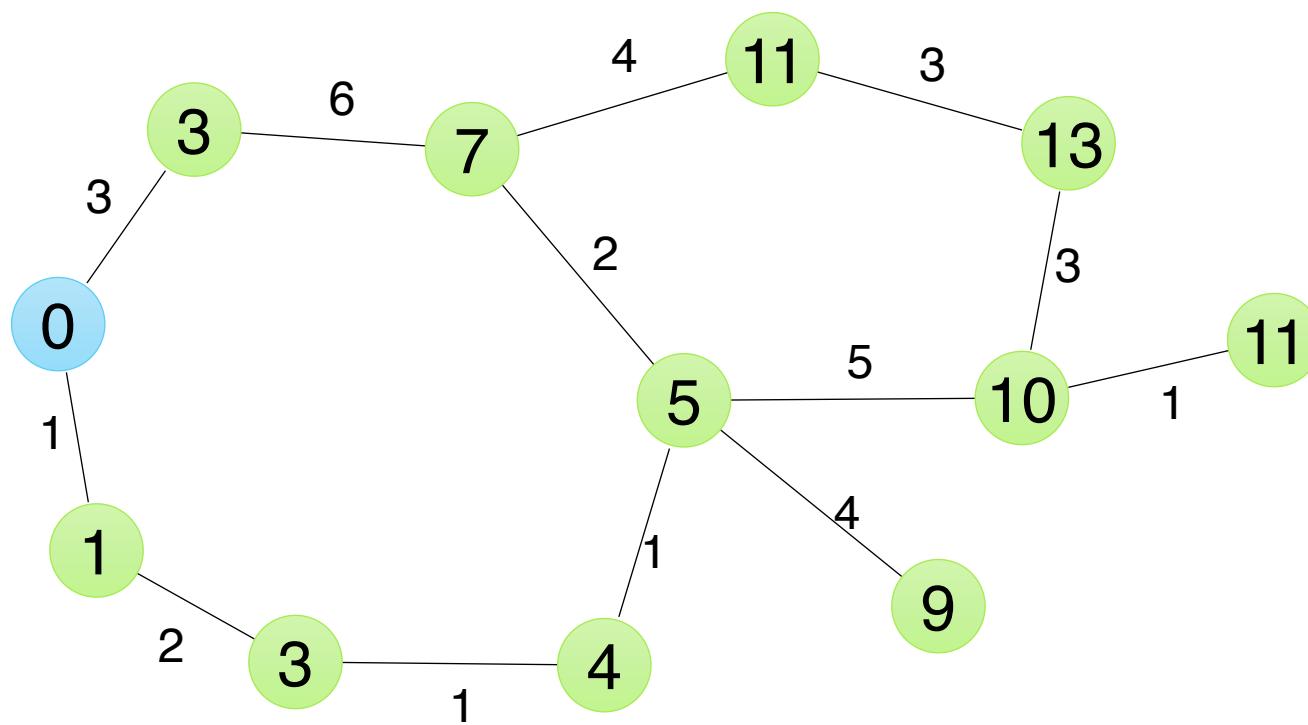
SSSP on Weighted Graphs



Work out!



SSSP on Weighted Graphs





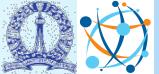
Complexity

- Using a **linked list** for queue, it takes $O(v^2 + e)$
- For each vertex,
 - we linearly search the linked list for smallest: $O(v)$
 - we check and update for each incident edge once: $O(d)$
- When a **min heap (priority queue)** with distance as priority key, total time is $O((v + e) \log v)$
 - $O(\log v)$ to insert or remove from priority queue
 - $O(v)$ *remove min* operations
 - $O(e)$ *change d[] value* operations (insert/update)
- When e is $O(v^2)$ [*highly connected, small diameter*], using a min heap is worse than using a linear list
- When a **Fibonacci heap** is used, the total time is $O(e + v \log v)$



Tasks

- Self study
 - Read: Graphs and graph algorithms (online sources)



Questions?